

ECSE 446/546

IMAGE SYNTHESIS

Monte Carlo Integration I



Prof. Derek Nowrouzezahrai

derek@cim.mcgill.ca

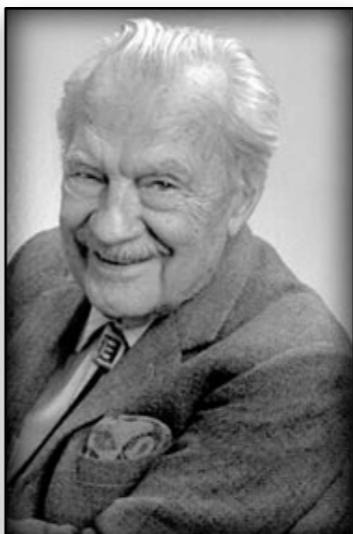
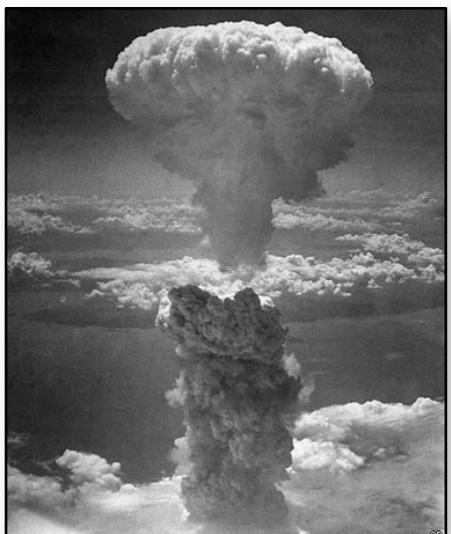
(slides in part by W. Jarosz)

Today

- Probability review
- Monte Carlo integration
- Random sampling

Monte Carlo History

- Use random numbers to solve numerical problems
- Early use during development of atomic bomb
- Von Neumann, Ulam, Metropolis
- Named after the casino in Monte Carlo



Monte Carlo Methods

- Pros
 - Flexible
 - Easy to implement
 - Easily handles complex integrands
 - Efficient for high dimensional integrands
- Cons
 - Variance (noise)
 - Slow convergence* $O(1/\sqrt{N})$

Random Variables

- Random variable X
- Cumulative distribution function (CDF)

$$P(x) = \text{Prob}\{X \leq x\} \quad P(x) \in [0, 1]$$

- Probability density function (PDF)

$$p(x) = \frac{dP(x)}{dx} \quad p(x) \geq 0 \quad \int p(x) = 1$$

- Therefore

$$\text{Prob}\{a \leq X \leq b\} = \int_a^b p(x) dx = P(b) - P(a)$$

Random Variables

- Uniform random variables:

$$p(x) = \frac{dP}{dx}(x) = \text{const}$$

- Canonical uniform random variable ξ

$$p(x) = \begin{cases} 1 & x \in [0, 1], \\ 0 & \text{otherwise.} \end{cases}$$

Expected Value & Variance

- Random variable $Y = f(X)$
- Expected Value: $E[Y] = \int_D f(x)p(x) dx$
- Variance: $V[Y] = E[(Y - E[Y])^2]$
- Properties:
 - $E[aY] = aE[Y]$
 - $E[Y_1 + Y_2] = E[Y_1] + E[Y_2]$
 - $V[aY] = a^2V[Y]$
- Therefore: $V[Y] = E[Y^2] - E[Y]^2$

Expected Value

- Expected value:

$$E[Y] = \int_D f(x)p(x)dx \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$

where the x_i are distributed according to $p(x_i)$

- Strong law of large numbers:

$$\text{Prob} \left\{ E[Y] = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{i=1}^N f(x_i) \right\} = 1$$

Monte Carlo Integration

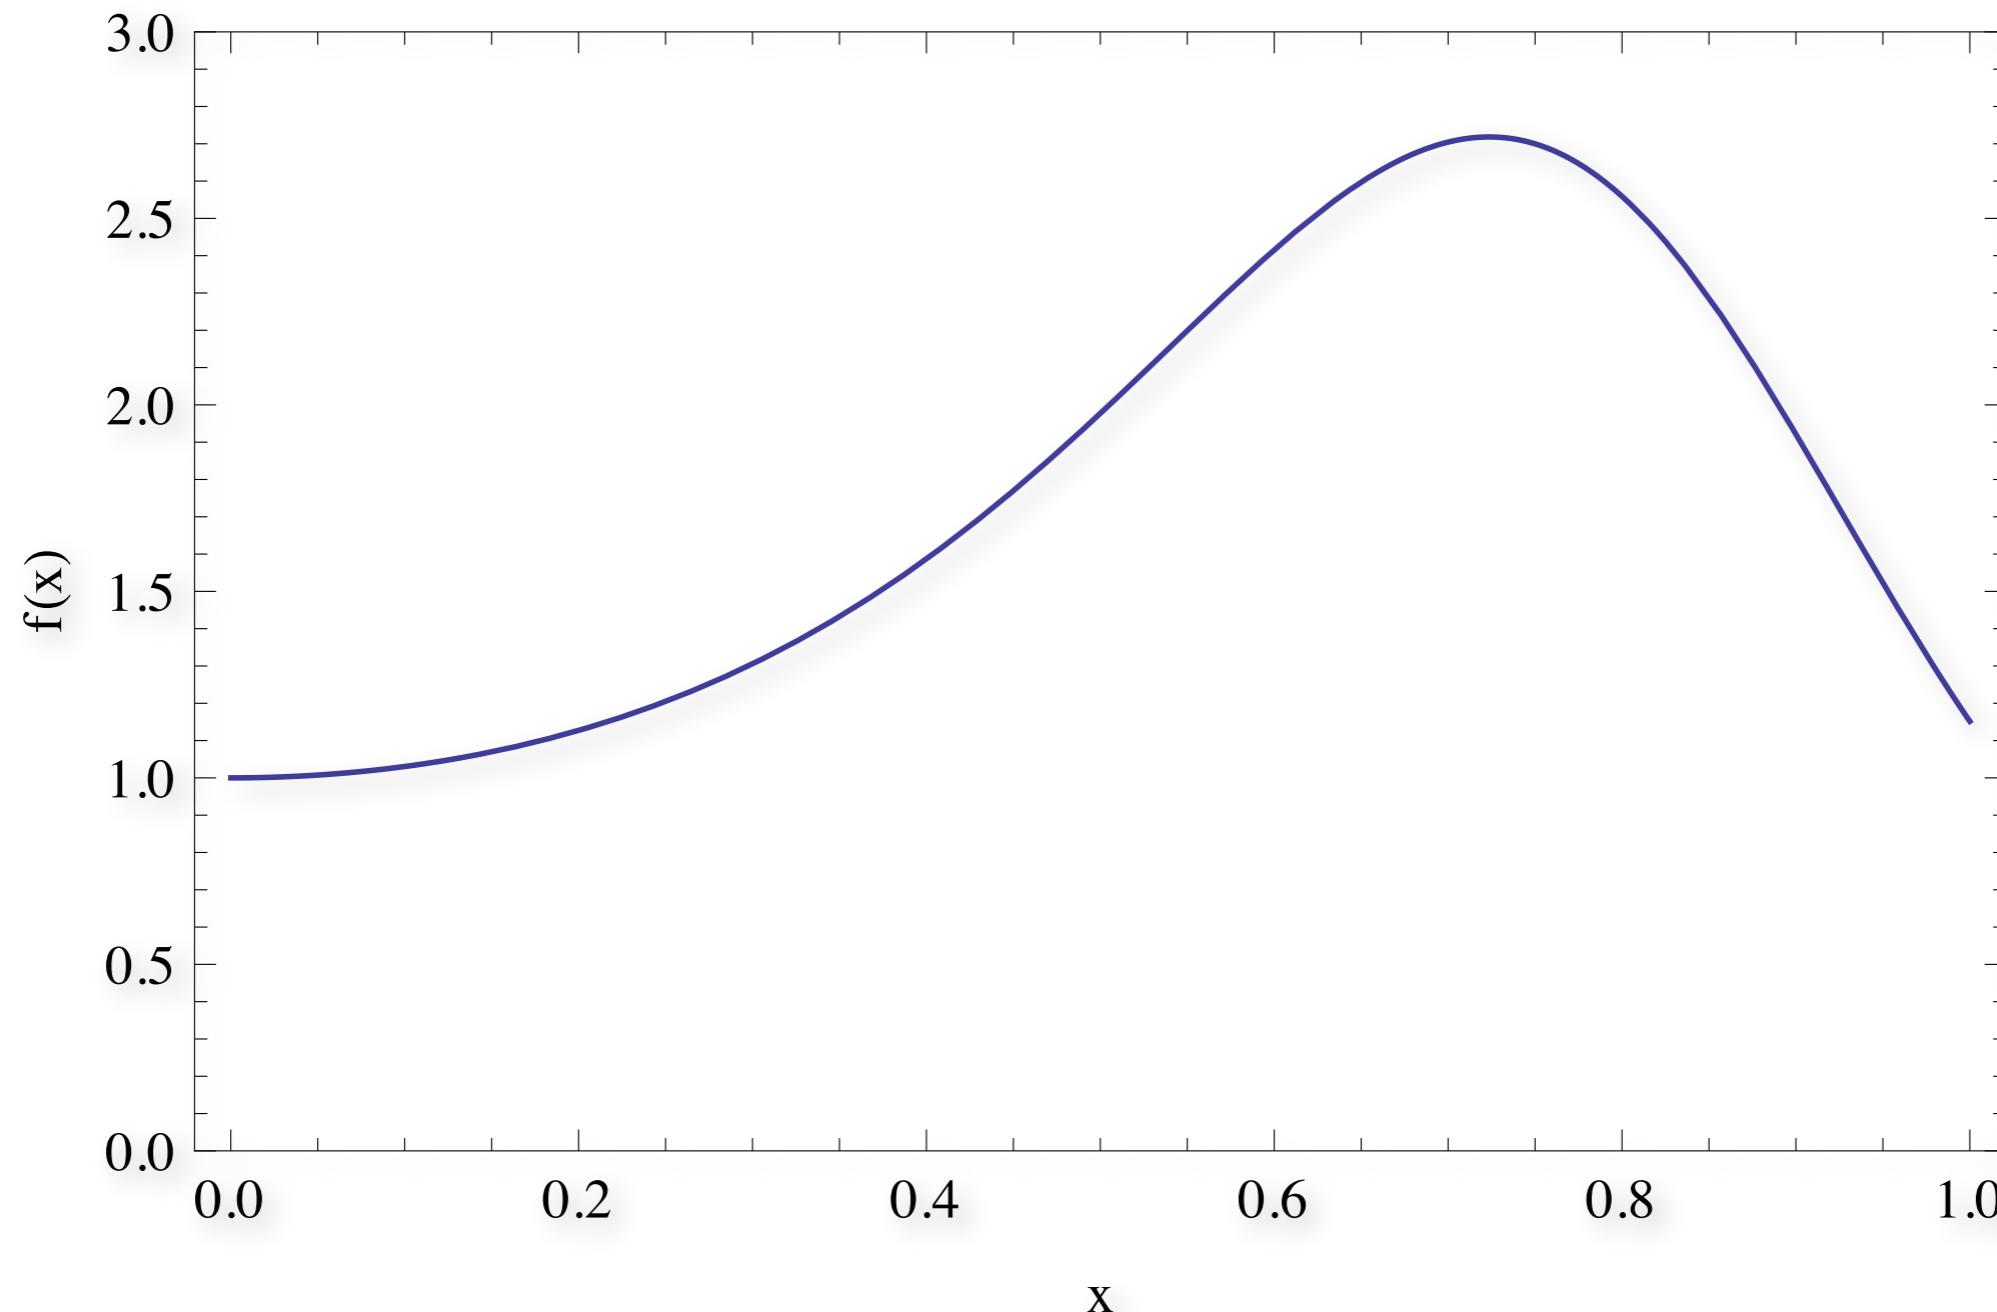
- We want to compute the value F of an integral

$$E[Y] = \int_D f(x)p(x)dx \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$

$$F = \int_D \left[\frac{f(x)}{p(x)} \right] p(x) dx \approx \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)} = F_N$$

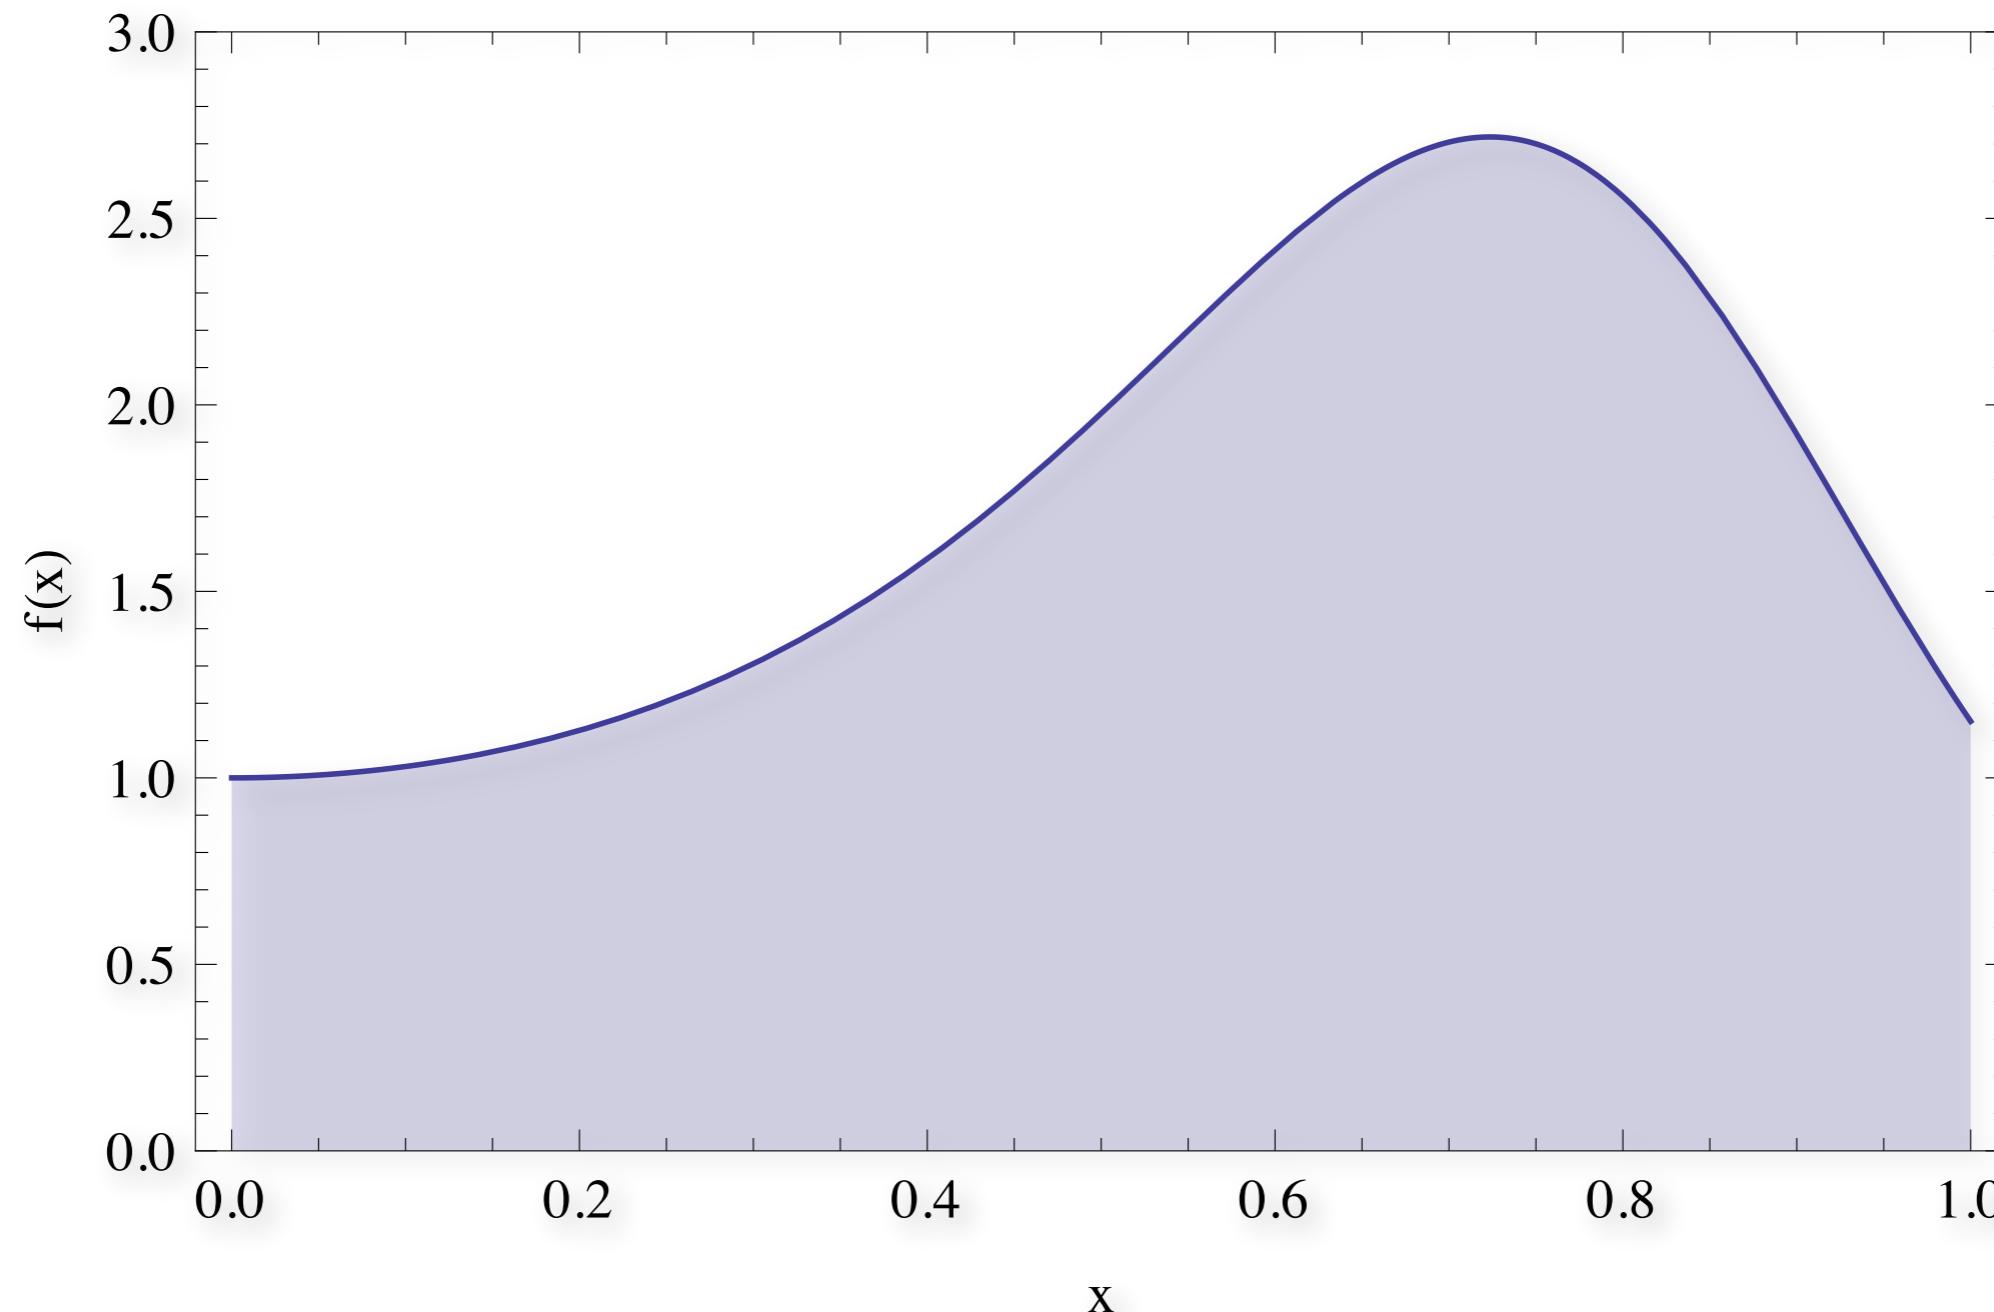
Monte Carlo Integration

$$f(x) = e^{\sin(3x^2)}$$



Monte Carlo Integration

$$F = \int_0^1 e^{\sin(3x^2)} dx$$



Monte Carlo Integration

$$F = \int_0^1 e^{\sin(3x^2)} dx \approx F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)} \Rightarrow \frac{1}{N} \sum_{i=1}^N f(x_i)$$

```
double integrate(int N)
{
    double x, sum=0.0;
    for (int i = 0; i < N; ++i) {
        x = randf();                                 $p(x_i) = 1$ 
        sum += exp(sin(3*x*x));
    }
    return sum / double(N);
}
```

Monte Carlo Integration

$$F = \int_a^b e^{\sin(3x^2)} dx \approx F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)}$$

```
double integrate(int N, double a, double b)
{
    double x, sum=0.0;
    for (int i = 0; i < N; ++i) {
        x = randf();
        sum += exp(sin(3*x*x));
    }
    return sum / double(N);
}
```

Monte Carlo Integration

$$F = \int_a^b e^{\sin(3x^2)} dx \approx F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)}$$

```
double integrate(int N, double a, double b)
{
    double x, sum=0.0;
    for (int i = 0; i < N; ++i) {
        x = a + randf()*(b-a);
        sum += exp(sin(3*x*x));
    }
    return sum / double(N);
}
```

$$p(x_i) = \frac{1}{b - a}$$

Monte Carlo Integration

$$F = \int_a^b e^{\sin(3x^2)} dx \approx F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)} \Rightarrow \frac{b-a}{N} \sum_{i=1}^N f(x_i)$$

```
double integrate(int N, double a, double b)
{
    double x, sum=0.0;
    for (int i = 0; i < N; ++i) {
        x = a + randf()*(b-a);
        sum += exp(sin(3*x*x));
    }
    return sum * (b-a) / double(N);
}
```

$$p(x_i) = \frac{1}{b-a}$$

Monte Carlo Integration

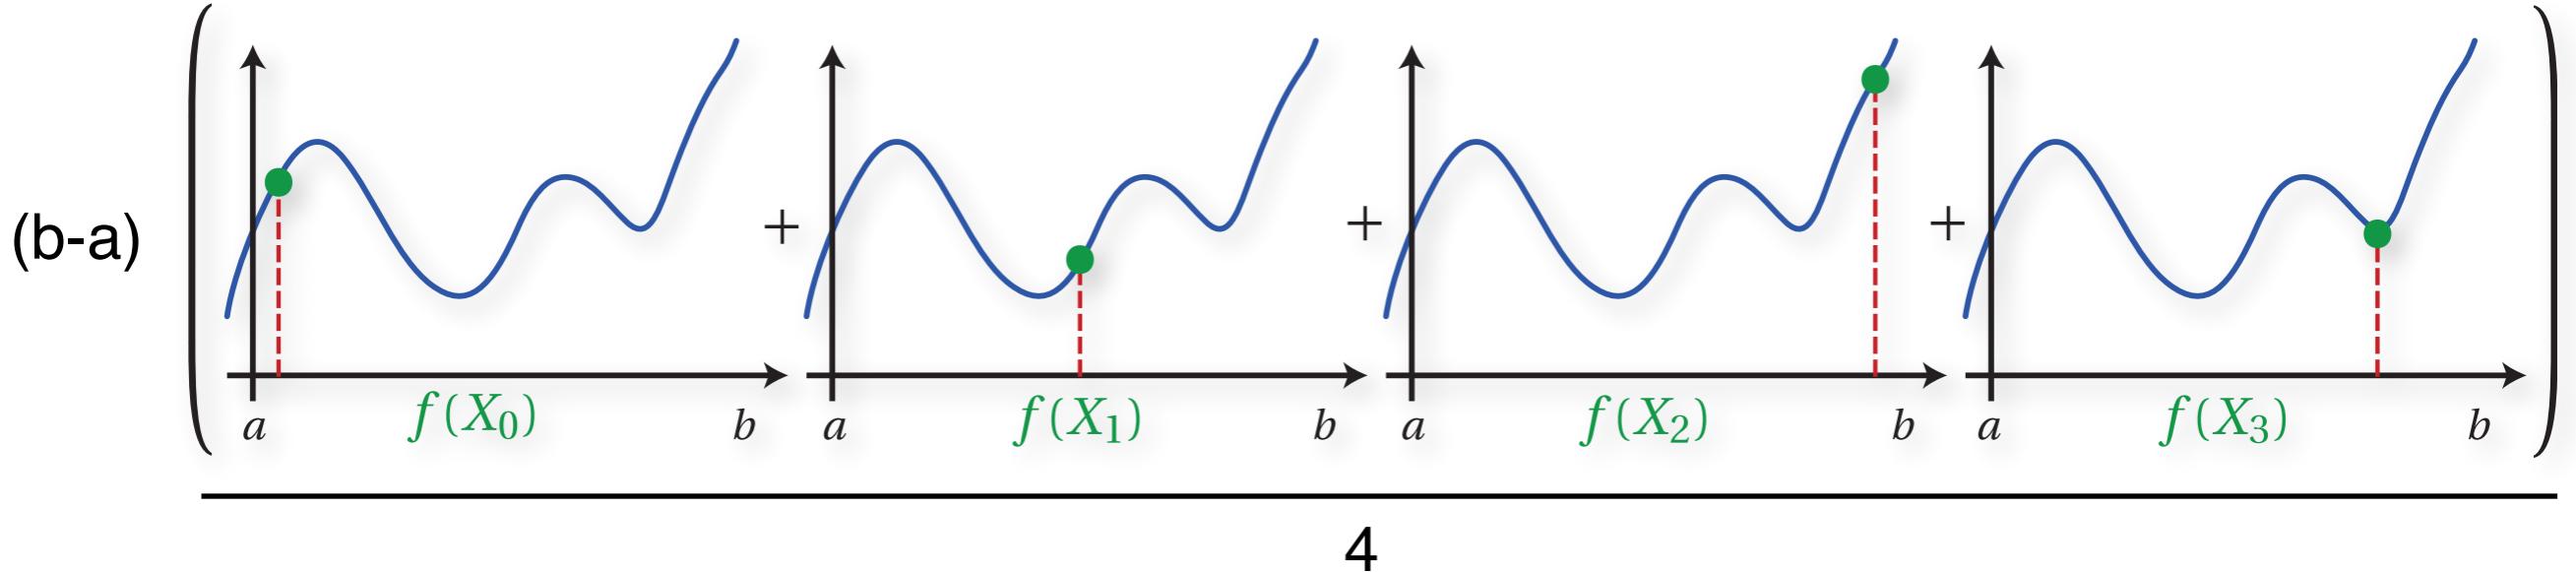
$$f(x) = e^{\sin(3x^2)}$$

N	F_N
1	2.75039
10	1.9893
100	1.79139
1000	1.75146
10000	1.77313
100000	1.77862

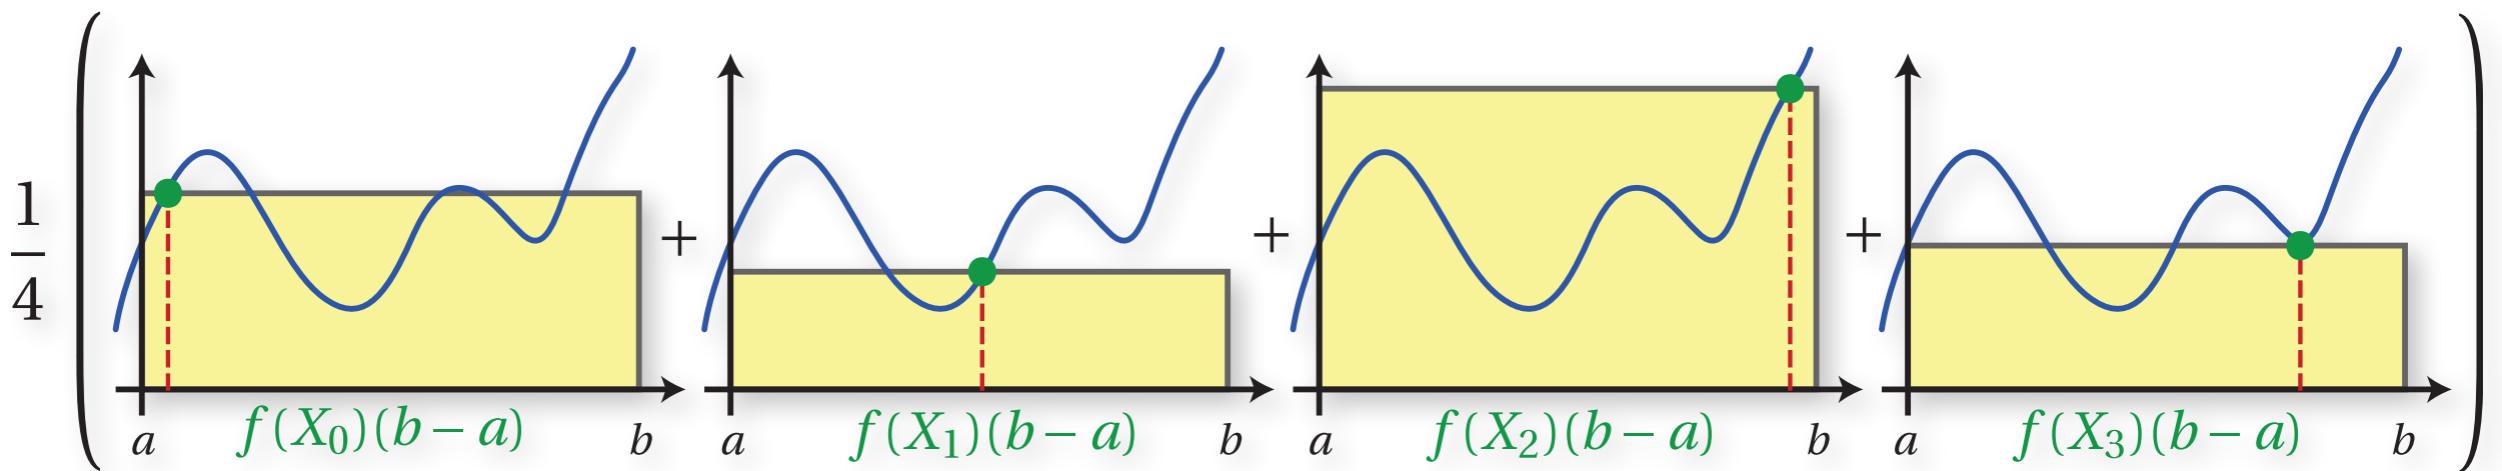
True value: 1.760977217585905...

Monte Carlo Visual Interpretation

Average function value or "height"



Area of a random box



Monte Carlo Integration Summary

- Goal: evaluate integral $\int_a^b f(x)dx$
- Random variable $X_i \sim p(x)$
- Monte Carlo Estimator $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(X_i)}{p(X_i)}$
- Expectation $E[F_N] = \int_a^b f(x)dx$

Monte Carlo Estimator

- *Unbiased estimator*
 - expected value equals integral
- Extension to higher dimensions straightforward
 - number of samples independent of dimension (cf. standard quadrature)
- Convergence* $O(1/\sqrt{N})$
 - reducing the error by a factor of 2 requires 4 times as many samples!

Today's Agenda

$$F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(X_i)}{p(X_i)}$$

- Main practical issues:
 - How to choose $p(x)$
 - How to generate x_i according to $p(x)$
- Case Study: **Ambient Occlusion**

$$L_r(\mathbf{x}, \vec{\omega}_r) = \int_{H^2} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_r) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\vec{\omega}_i$$

Visual Break



Sampling Random Variables

- Sampling the function domain:
 - Uniform unit interval (0,1)
 - Uniform interval (a,b)
 - Circle?
 - Sphere?
 - Hemisphere?
 - More complex domains?

Example: uniformly sampling a disk

- Uniform probability density on a unit disk

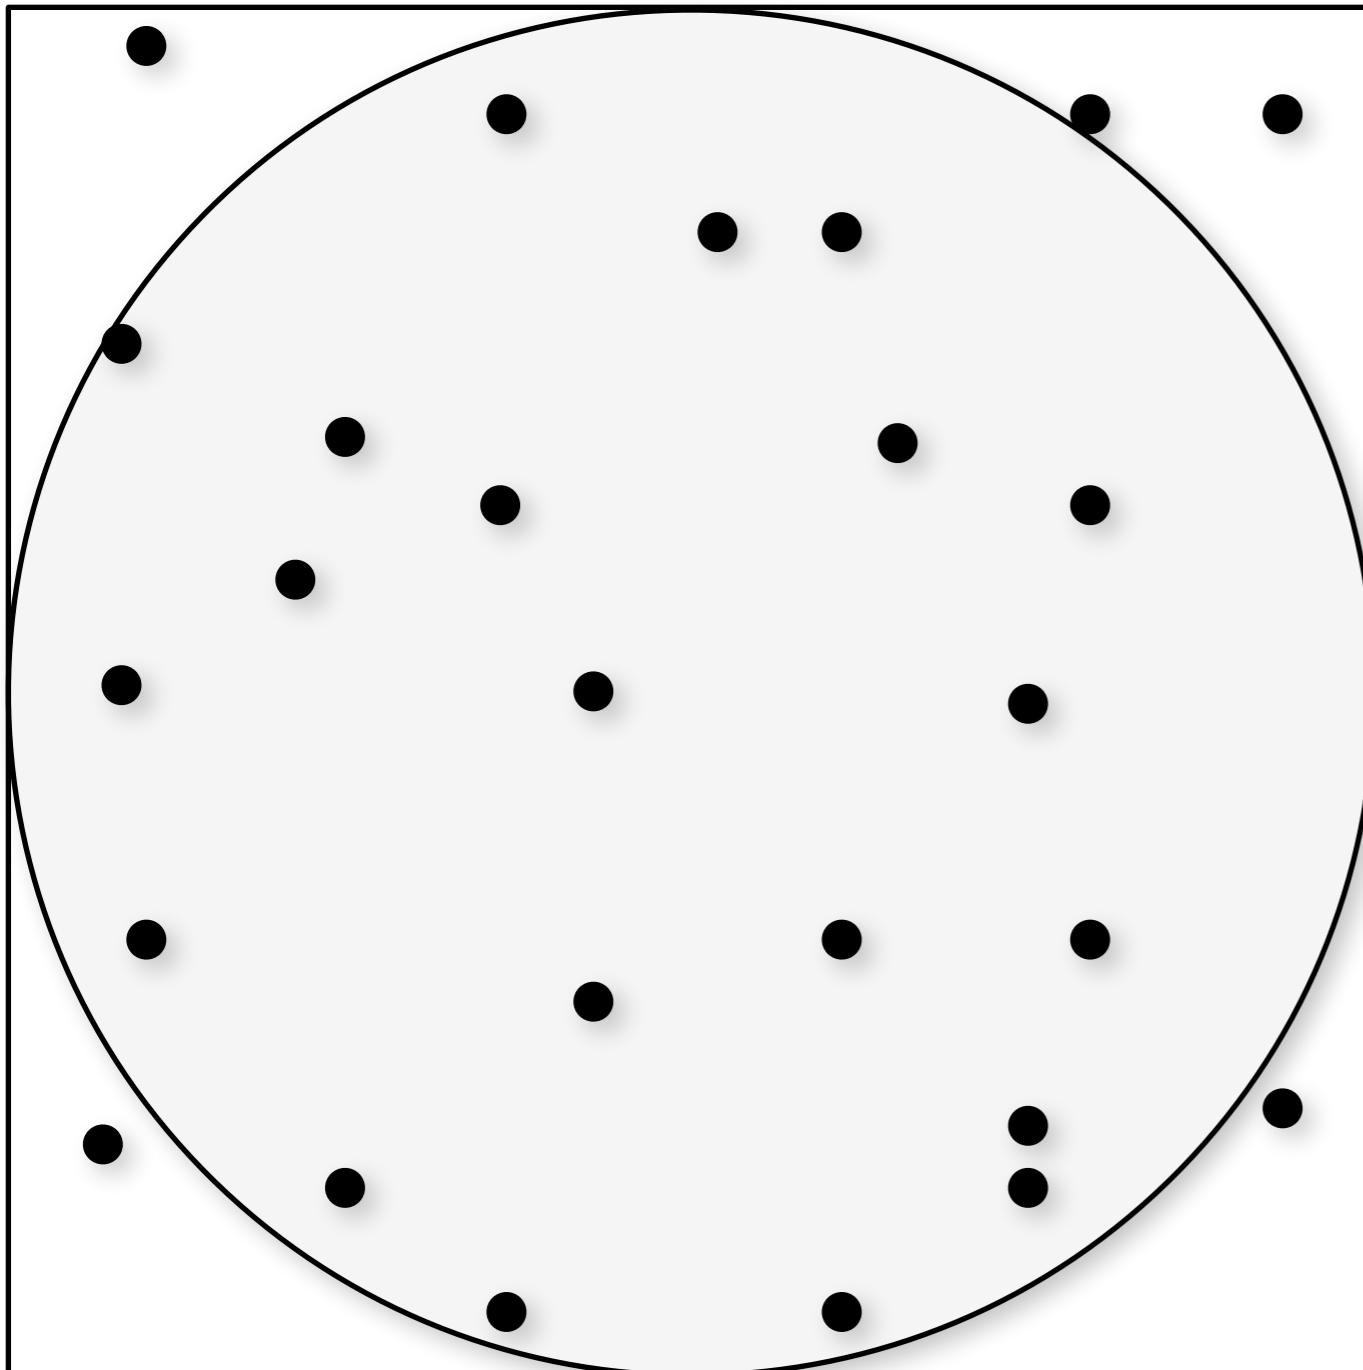
$$p(x, y) = \begin{cases} \frac{1}{\pi} & x^2 + y^2 < 1 \\ 0 & \text{otherwise} \end{cases}$$

- Goal: draw samples X_i, Y_i that are distributed as:

$$(X_i, Y_i) \sim p(x, y)$$

- Problem: pseudo-random number generators only allow us to draw samples from a canonical uniform distribution

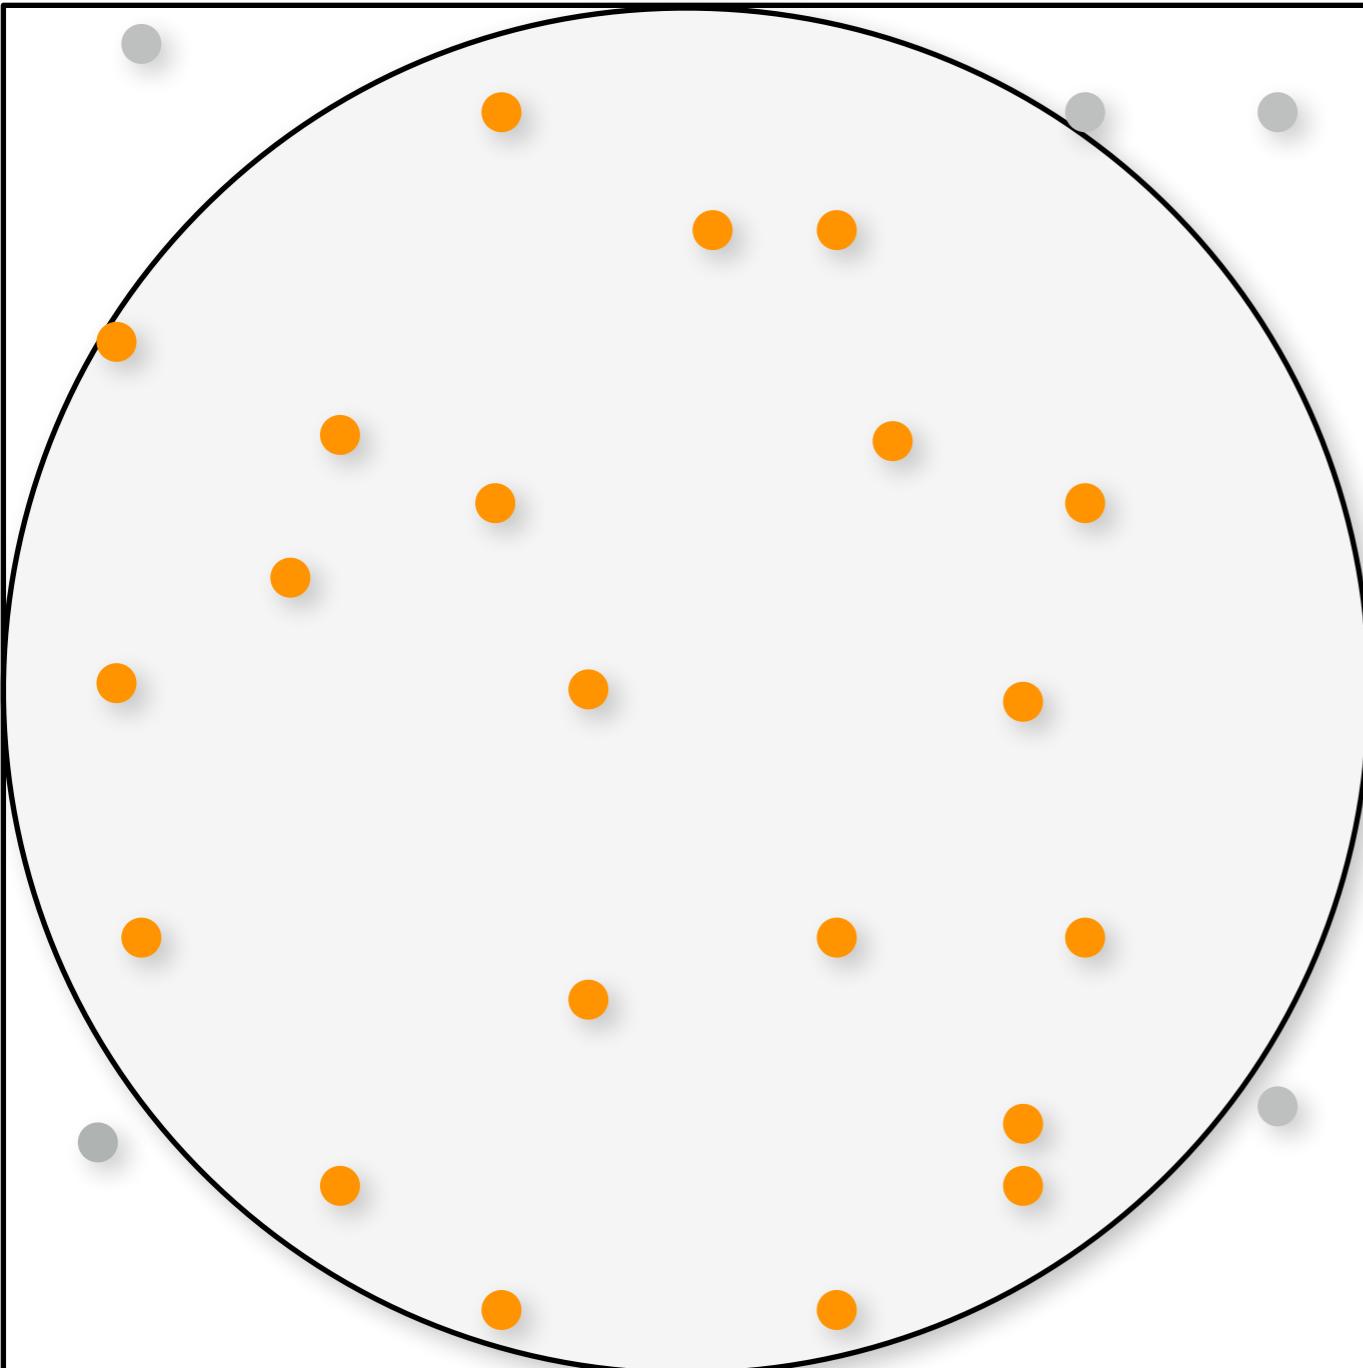
Rejection Sampling a Disk



```
Vector2 v;
```

```
v.x = 1-2*randf();  
v.y = 1-2*randf();
```

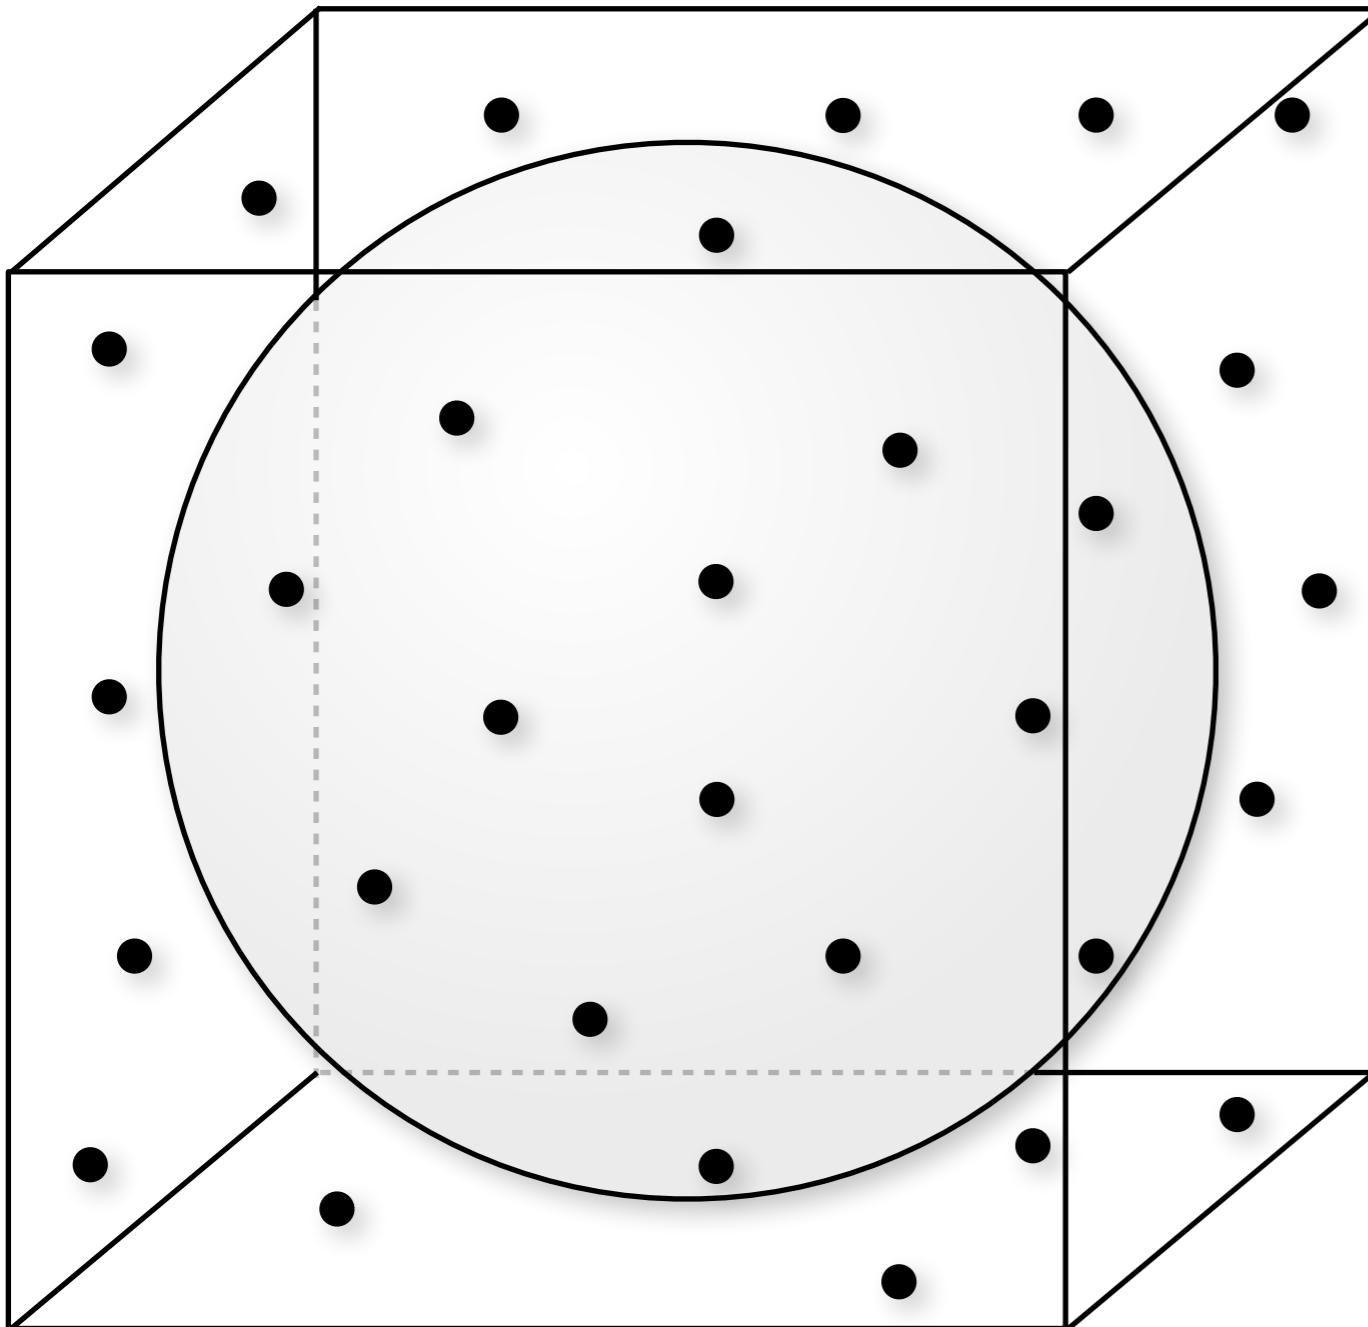
Rejection Sampling a Disk



```
Vector2 v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
} while(v.length2() > 1)
```

- Similar technique for sampling a sphere

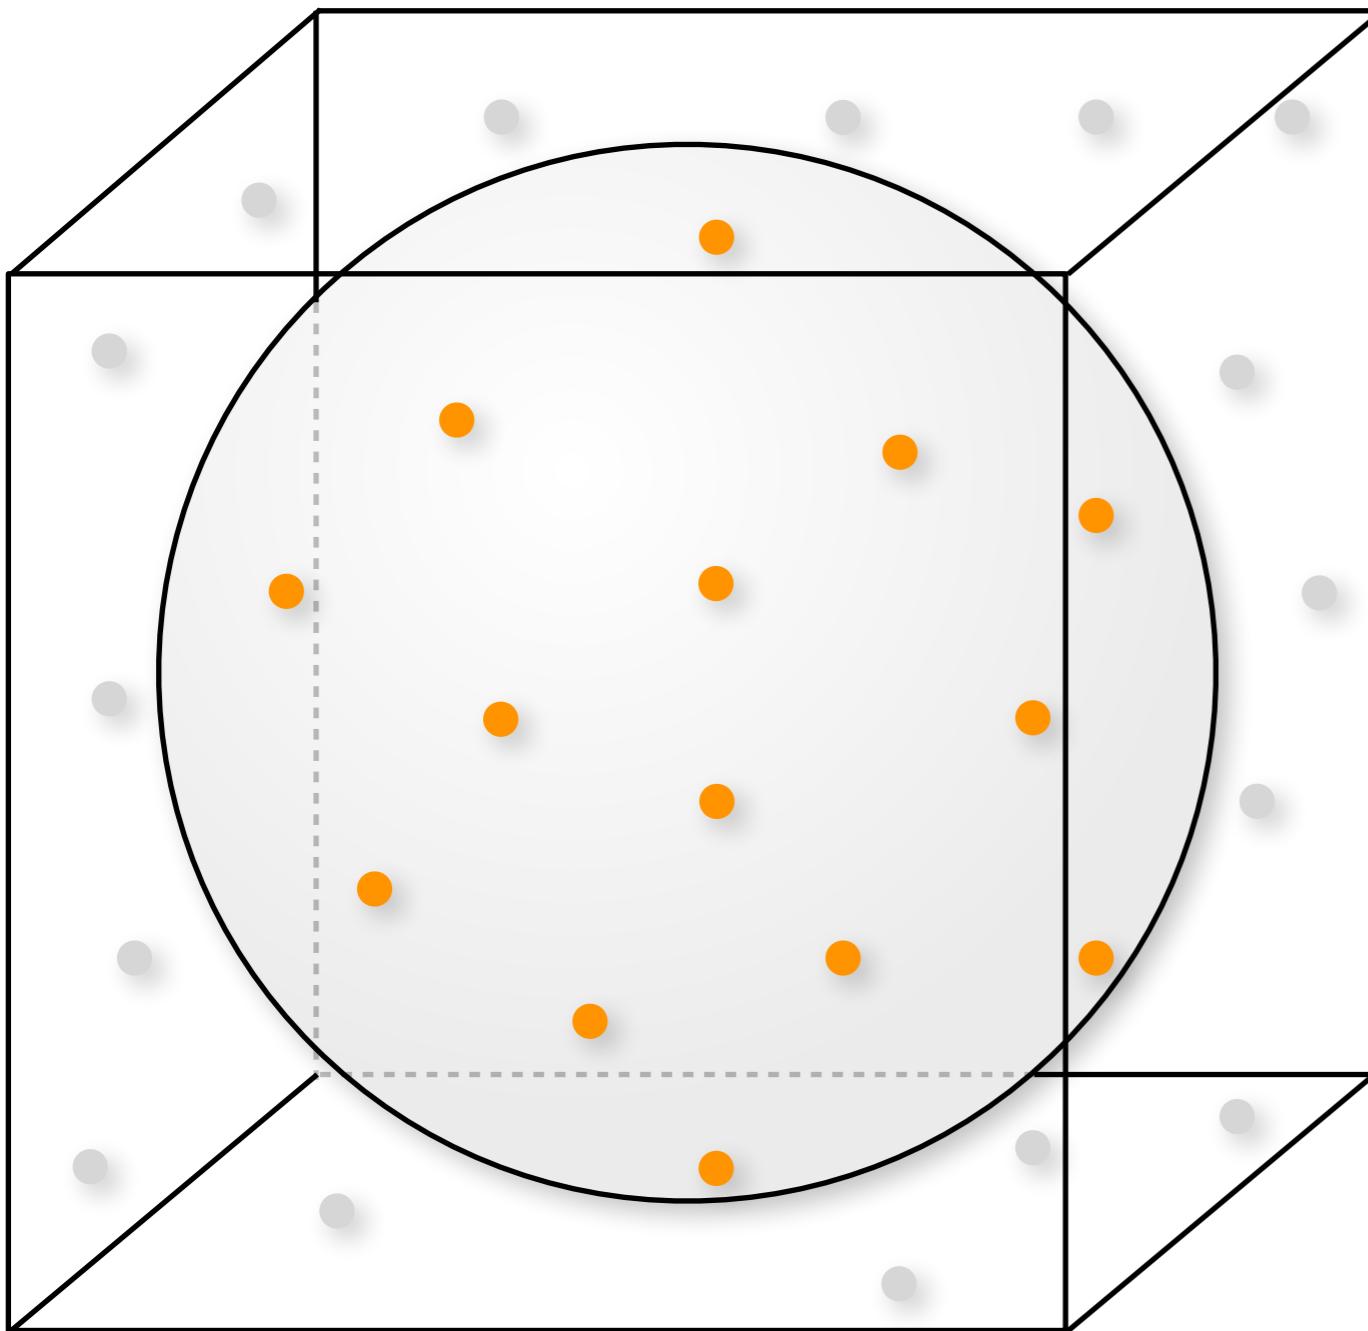
Rejection Sampling a Sphere



```
Vector3D v;
```

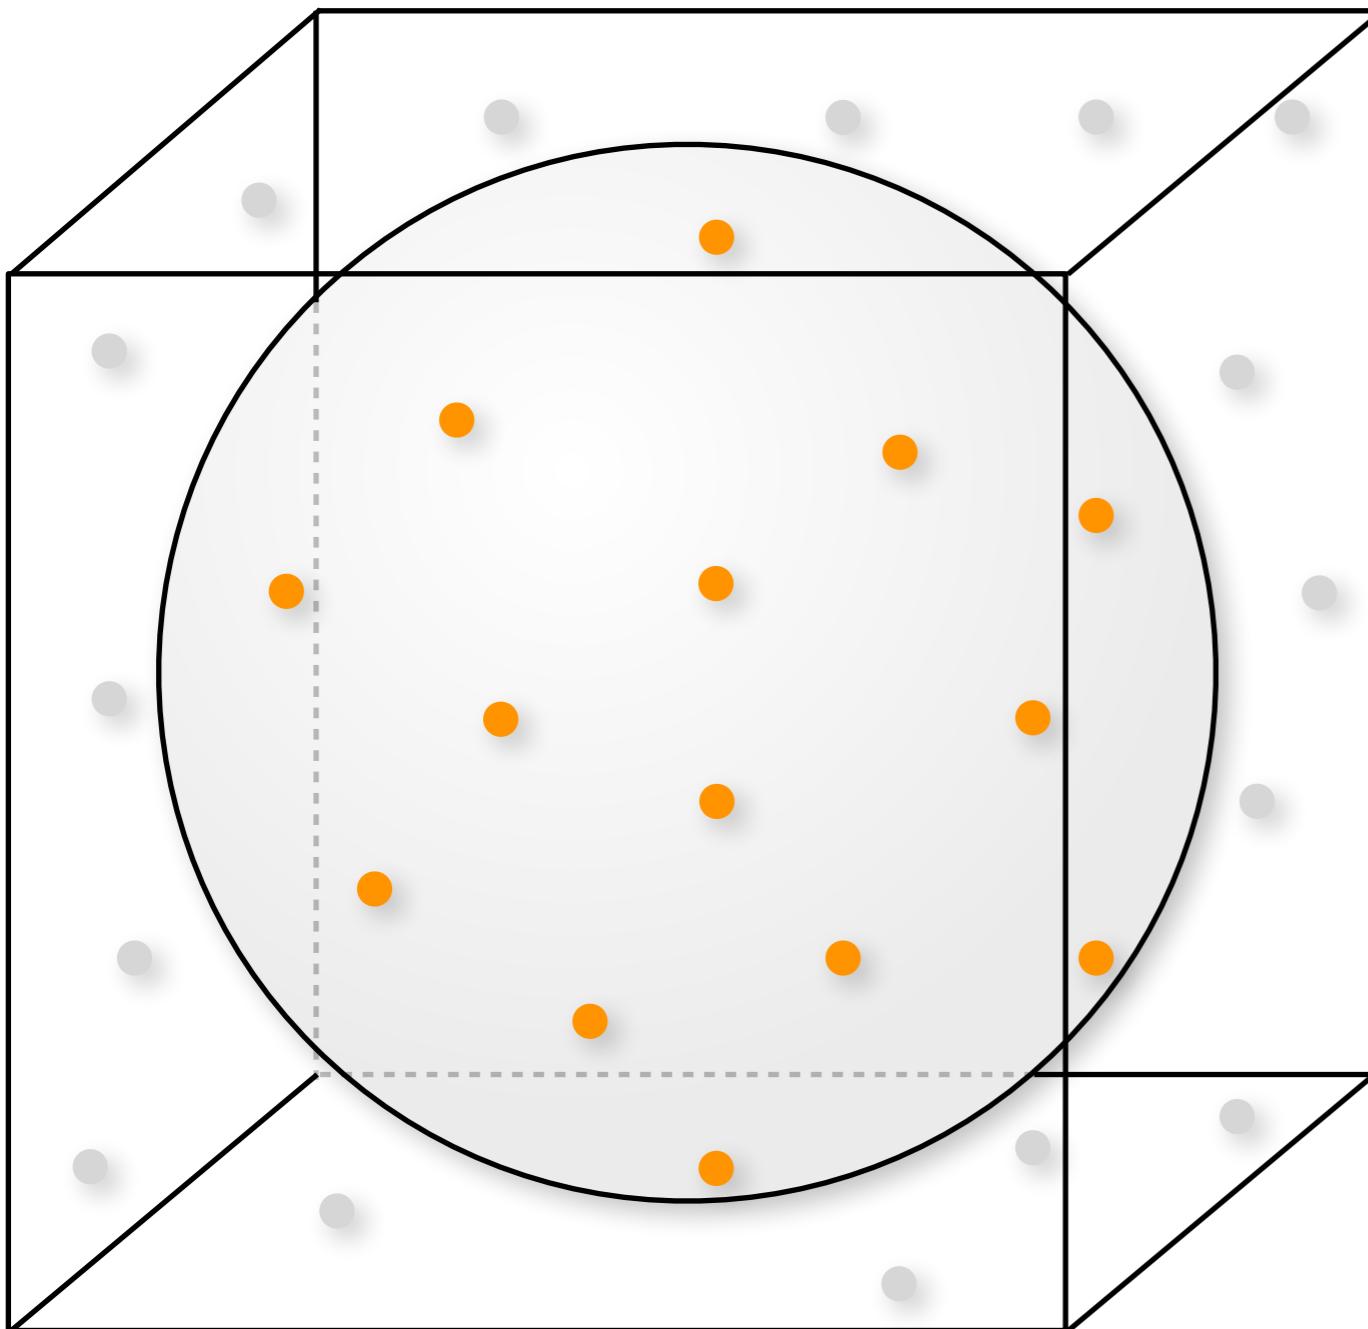
```
v.x = 1-2*randf();  
v.y = 1-2*randf();  
v.z = 1-2*randf();
```

Rejection Sampling a Sphere



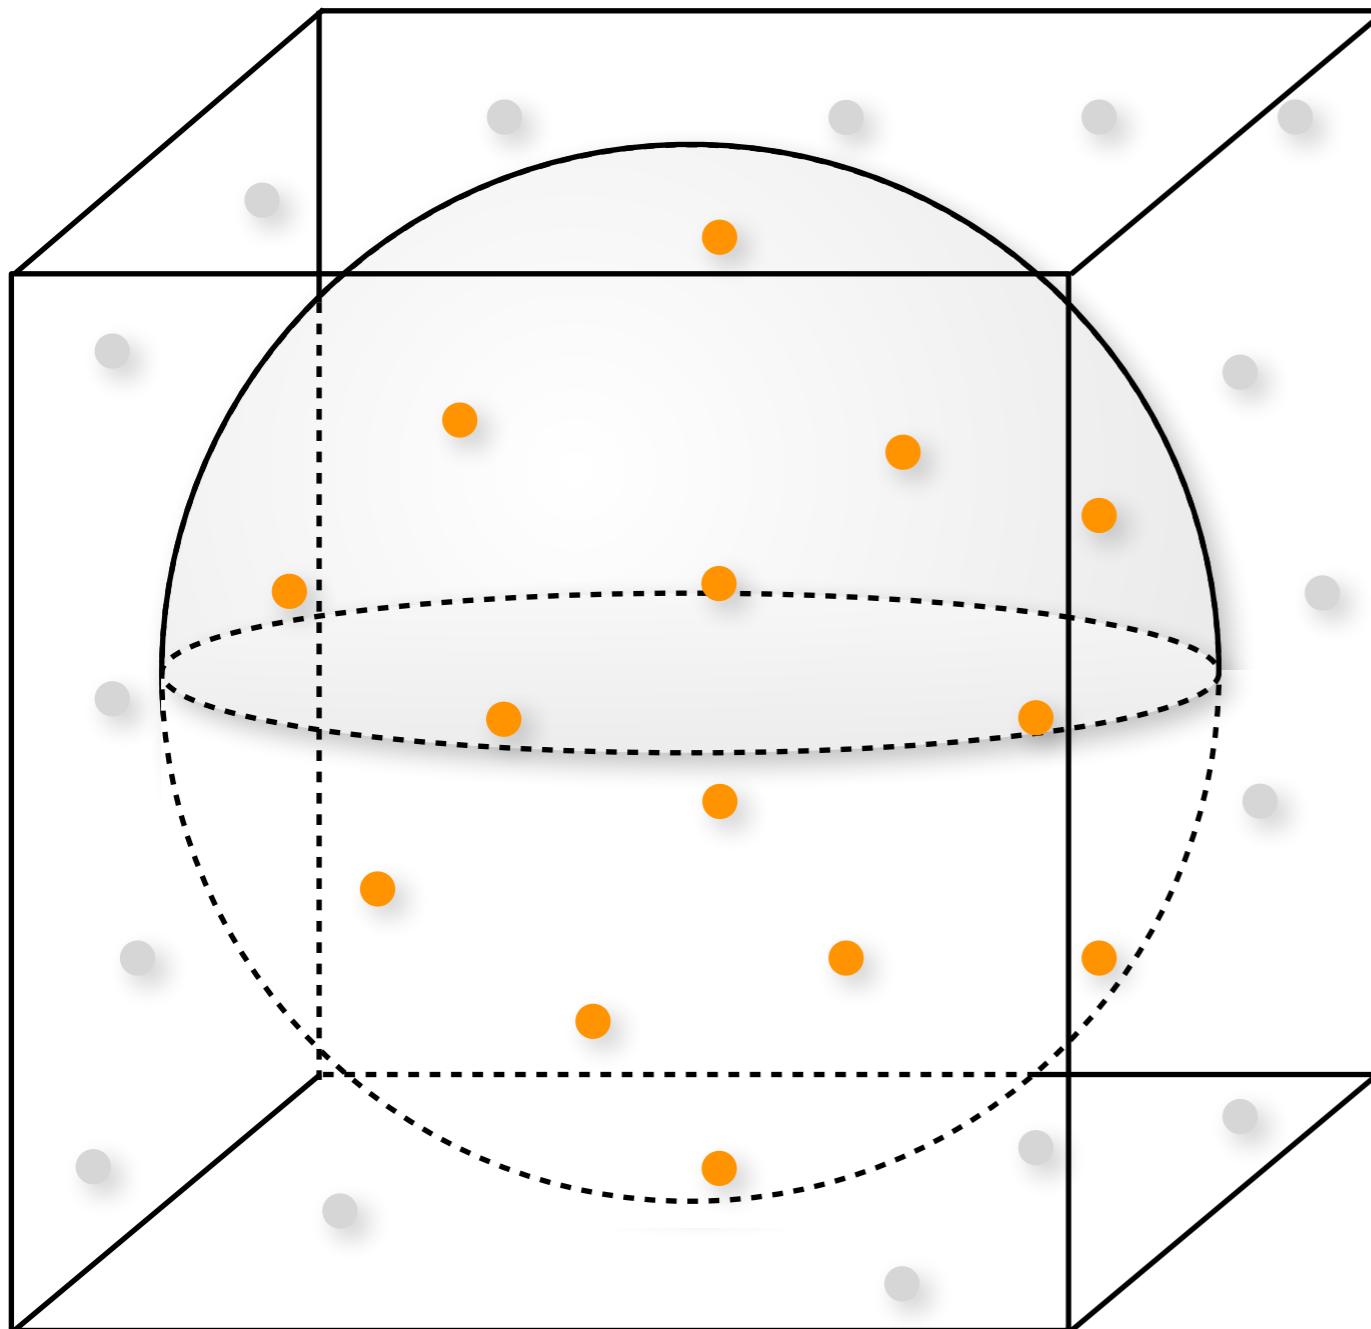
```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1)
```

Rejection Sampling a Sphere



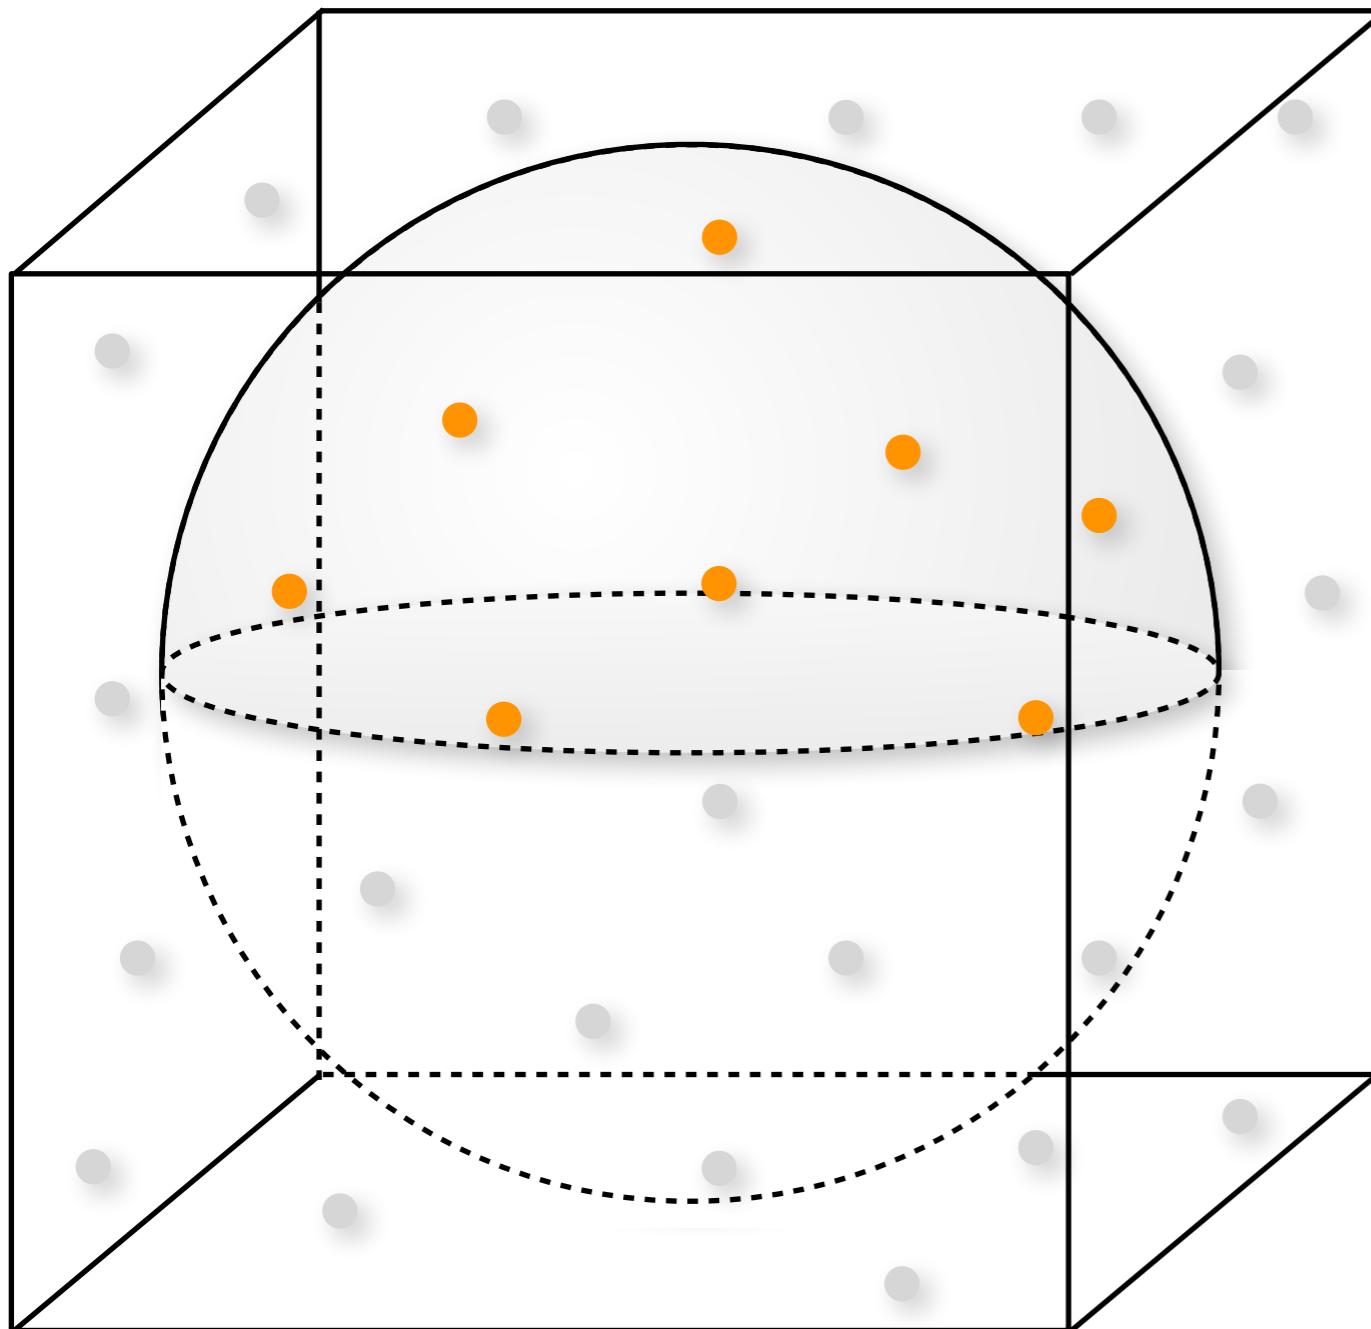
```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1)  
// Project onto sphere  
v /= v.length();
```

Rejection Sampling a Hemisphere



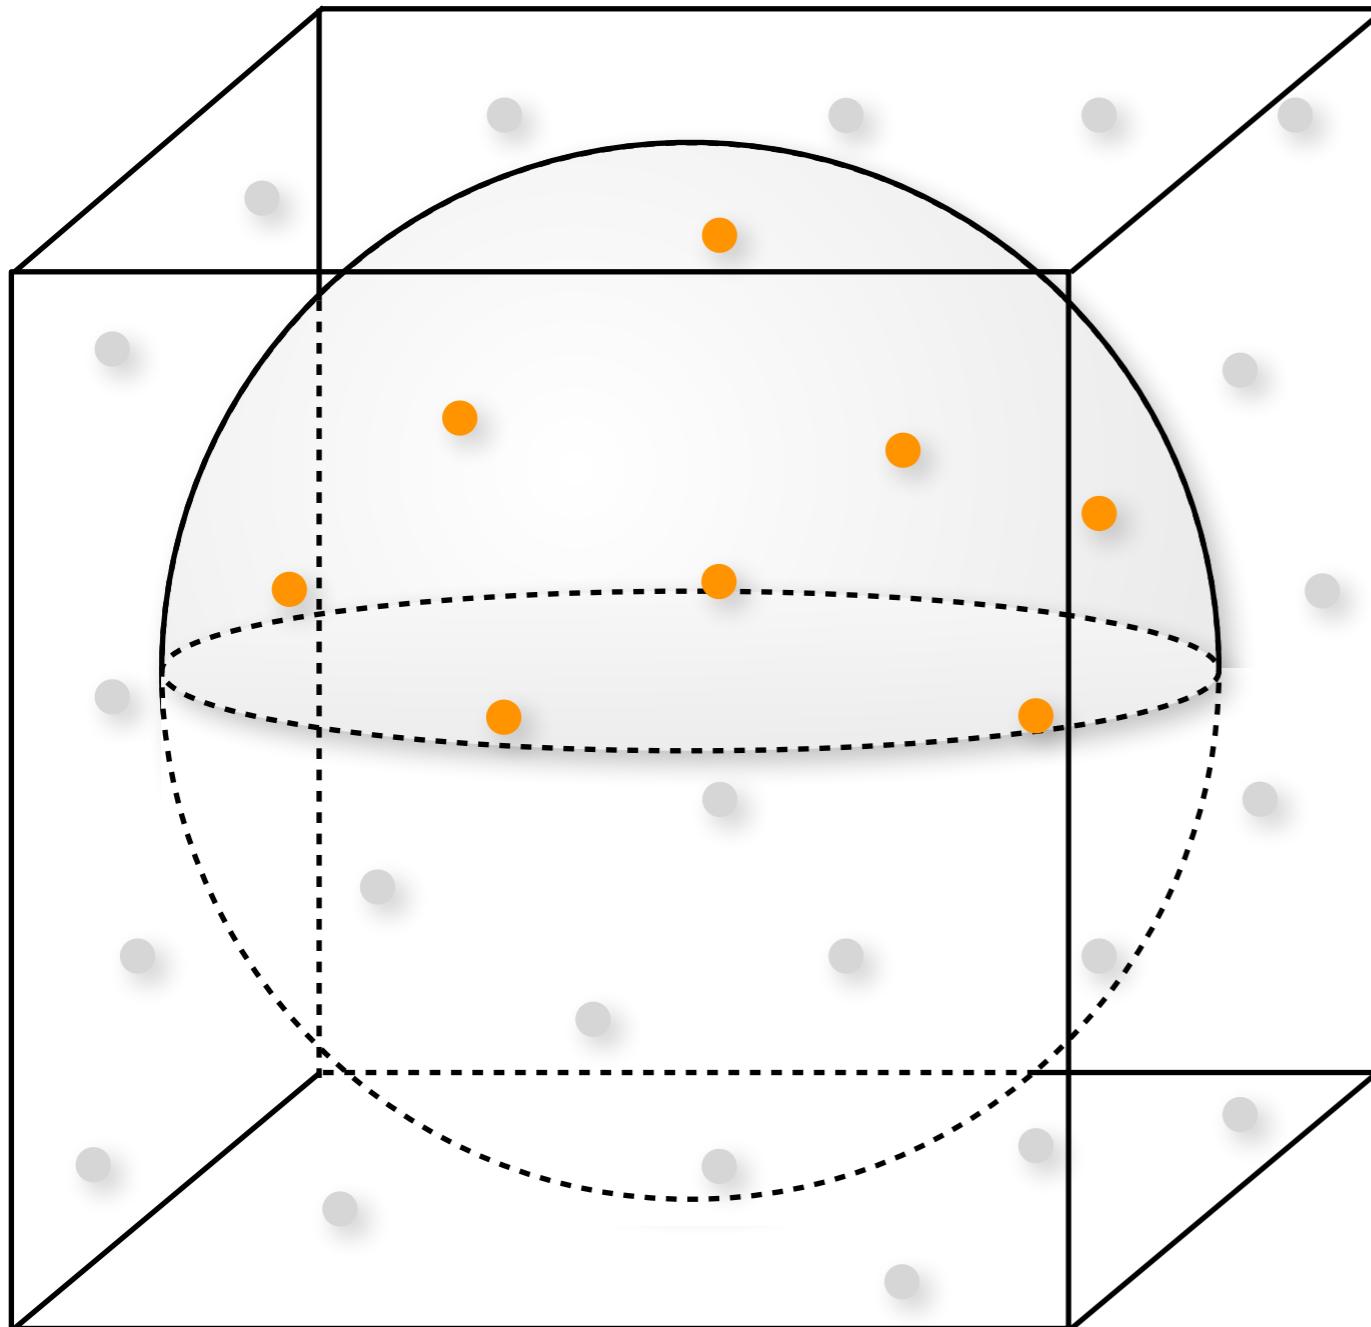
```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1)
```

Rejection Sampling a Hemisphere



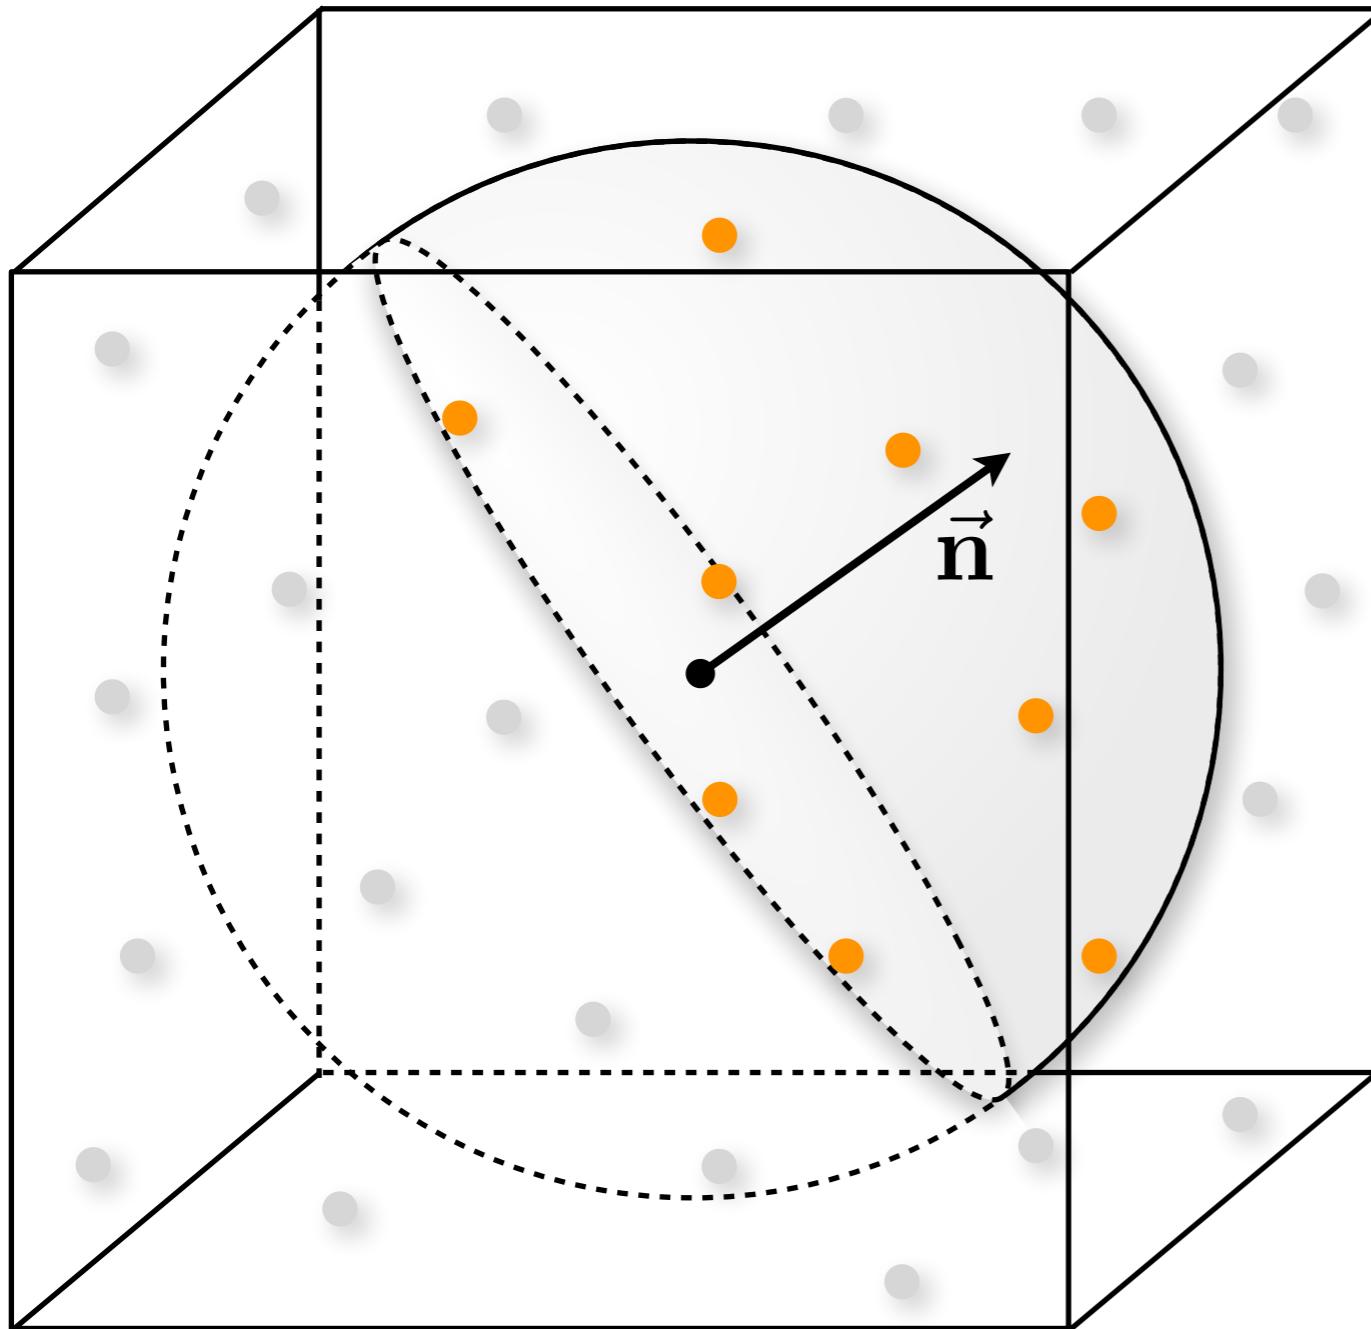
```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1 ||  
       v.z < 0)
```

Rejection Sampling a Hemisphere



```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1 ||  
       v.z < 0)  
  
// Project onto hemisphere  
v /= v.length();
```

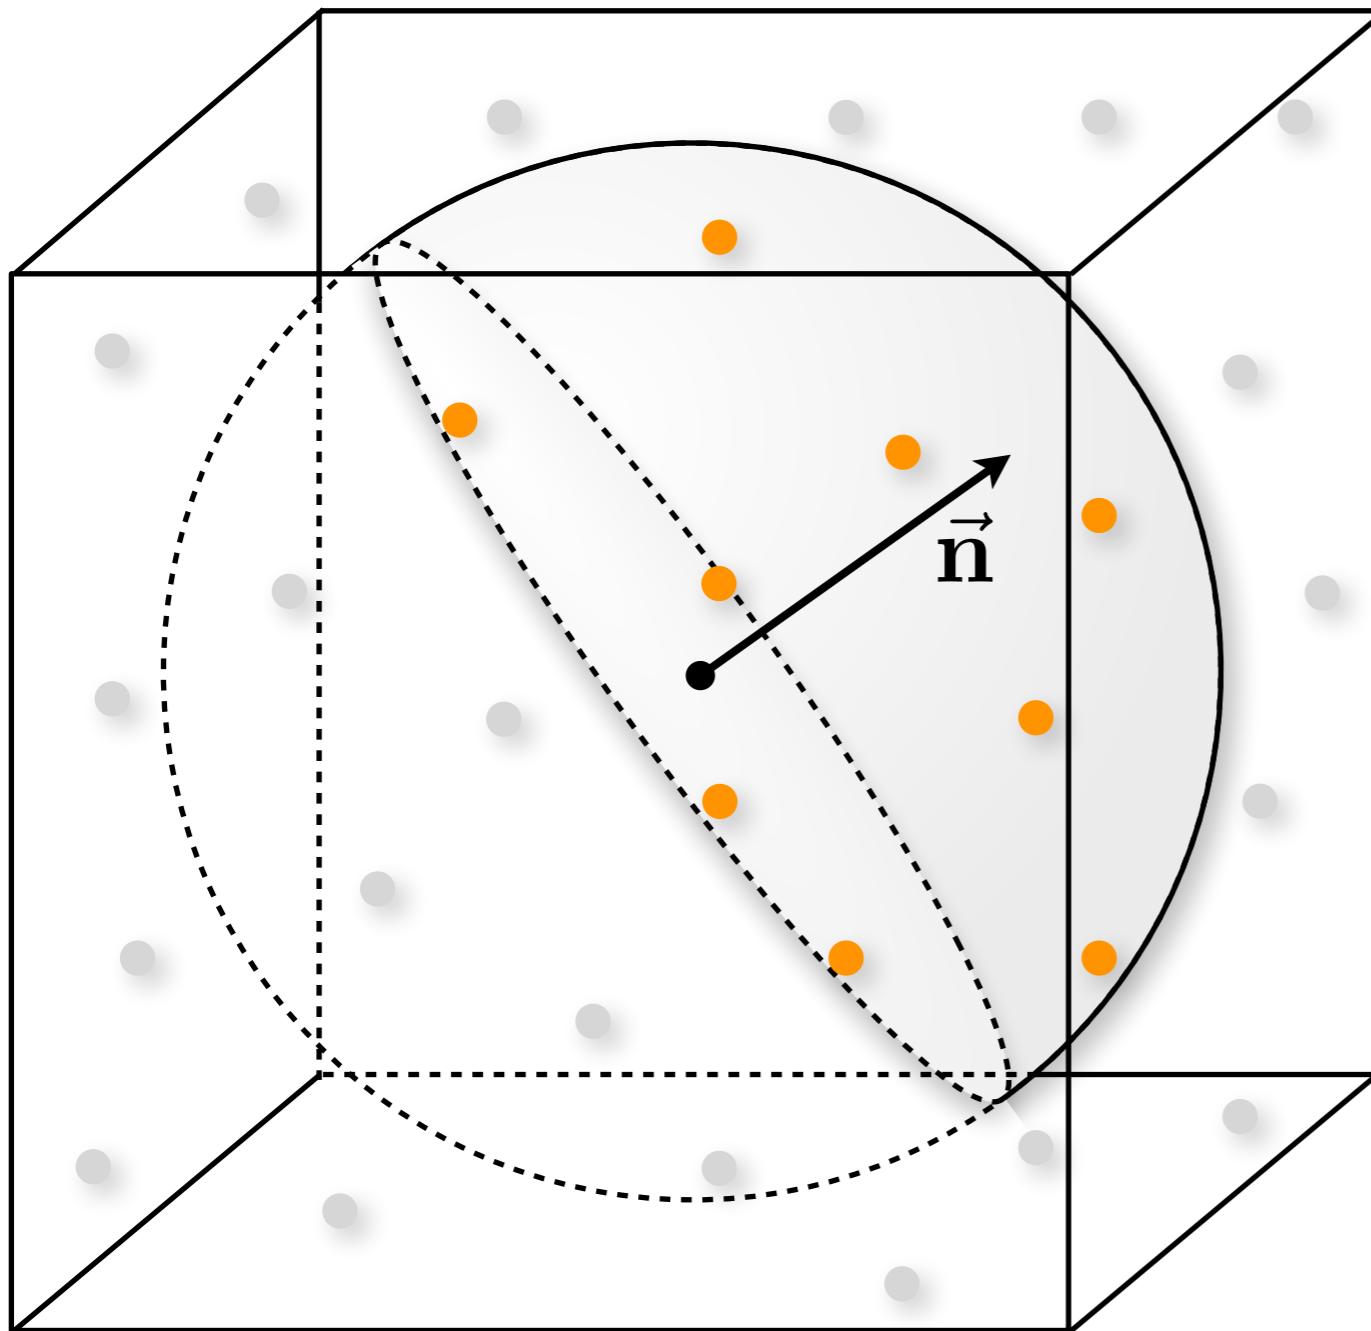
Rejection Sampling a Hemisphere



```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1 ||  
       v.z < 0)  
  
// Project onto hemisphere  
v /= v.length();
```

- Arbitrary orientation?

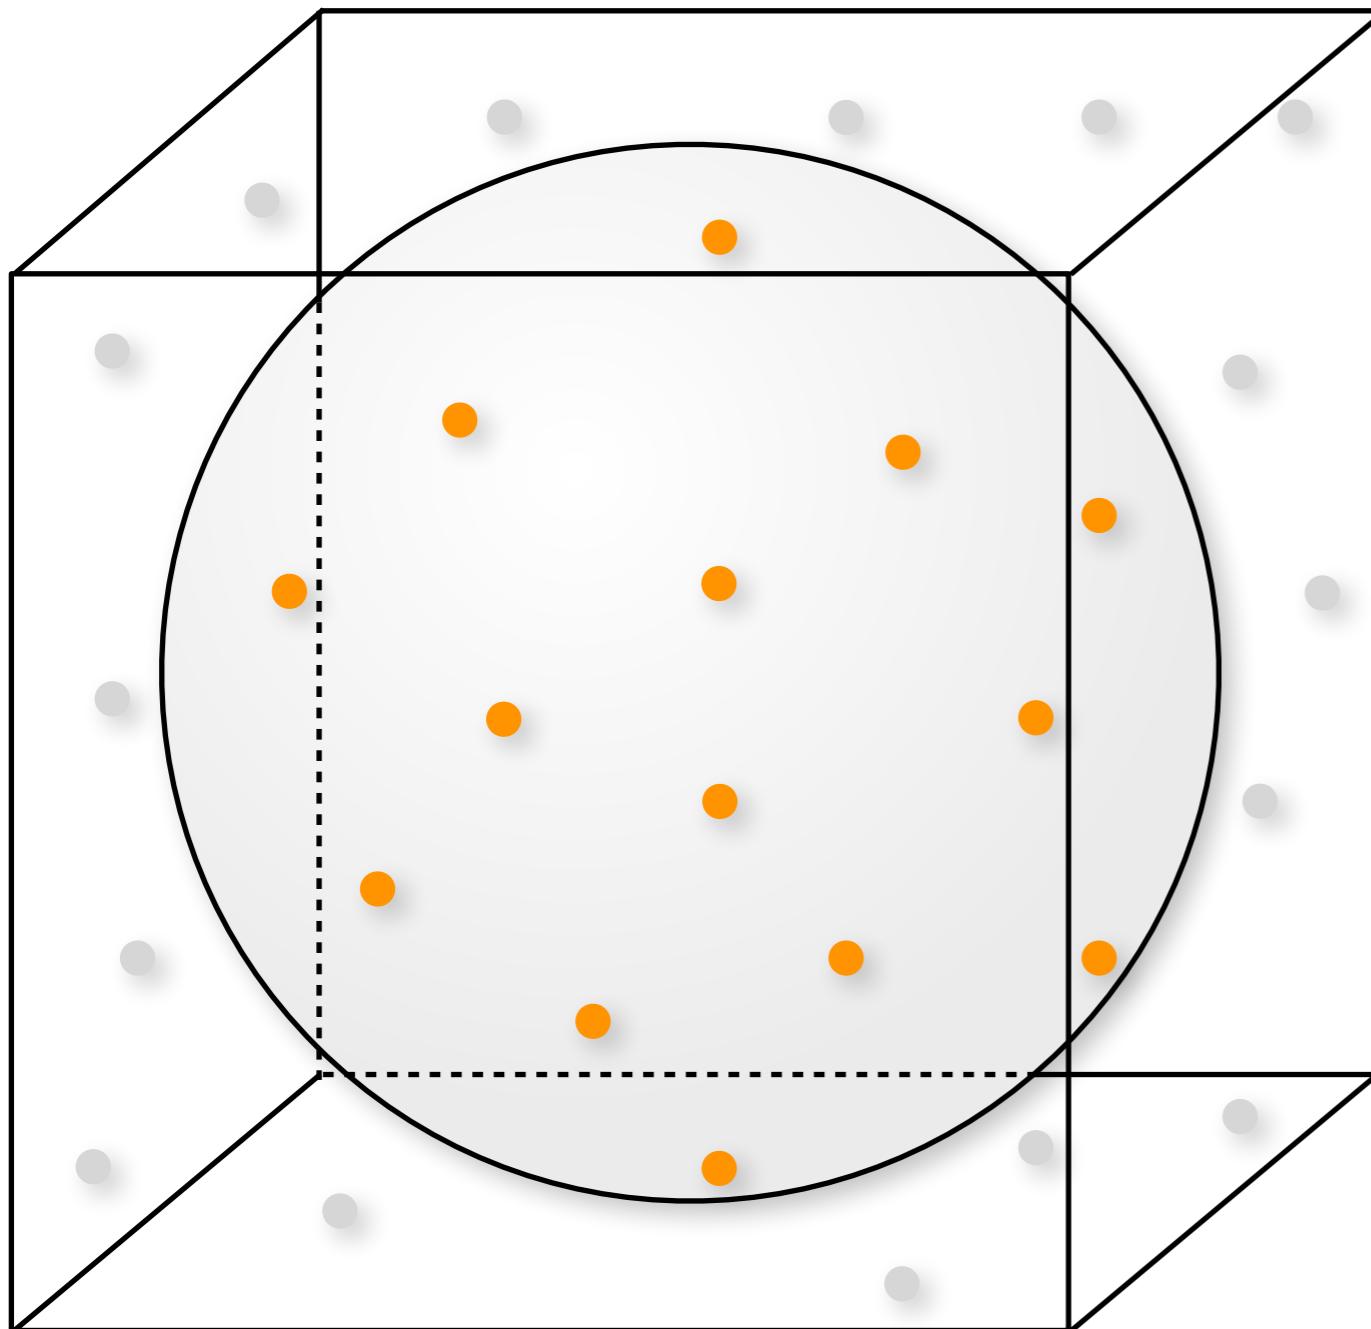
Rejection Sampling a Hemisphere



```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1 ||  
       dot(v,n) < 0)  
  
// Project onto hemisphere  
v /= v.length();
```

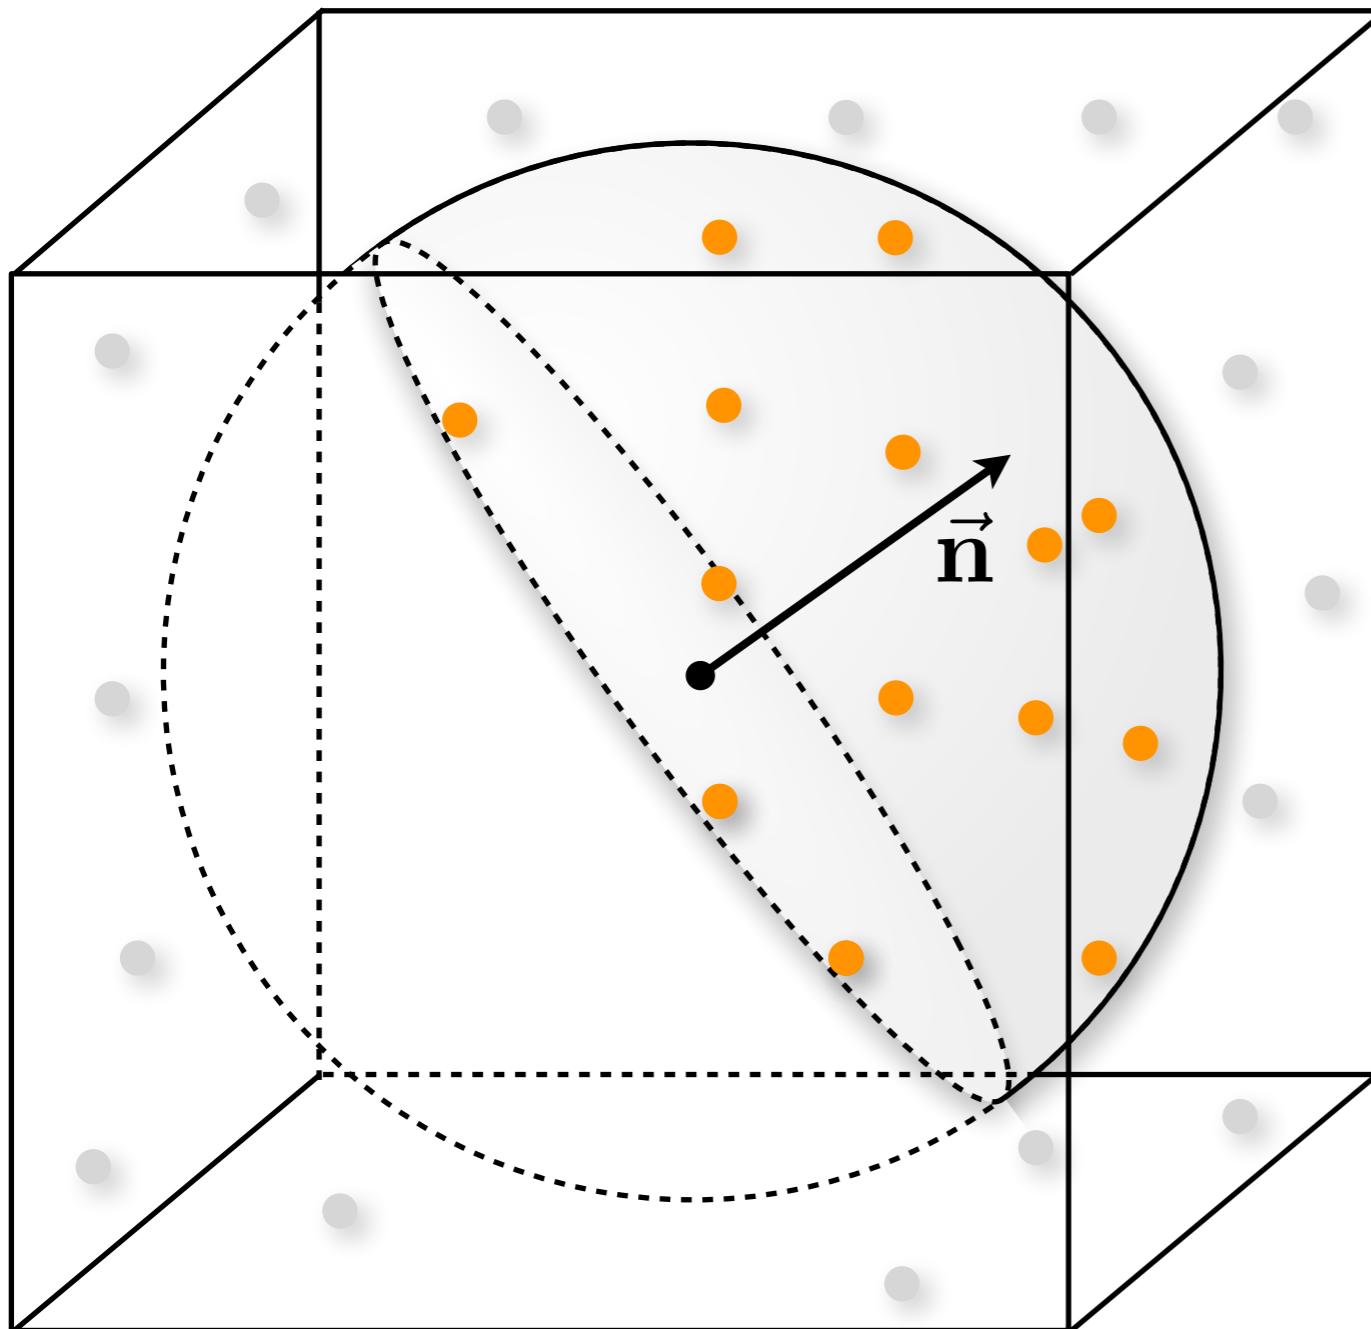
- Arbitrary orientation?
- **Inefficient!**

Rejection Sampling a Hemisphere



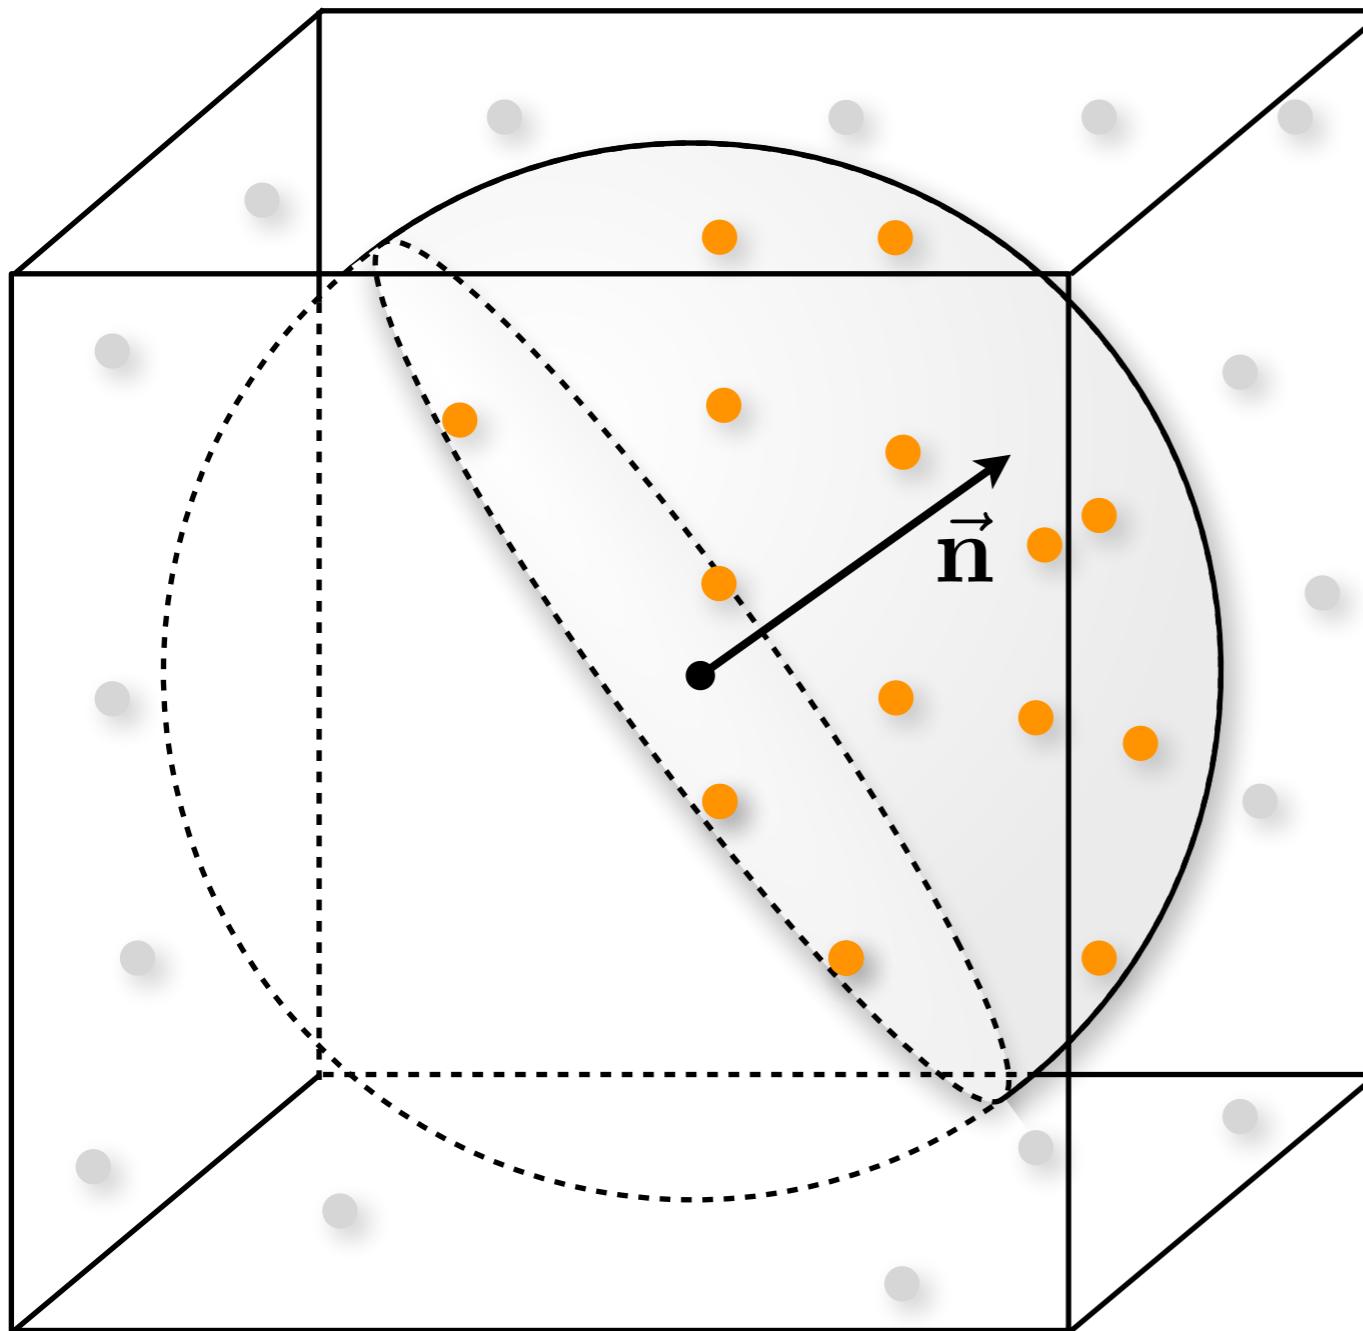
```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1)
```

Rejection Sampling a Hemisphere



```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1)  
// flip to proper hemisphere  
if (dot(v,n) < 0)  
    v = -v;
```

Rejection Sampling a Hemisphere

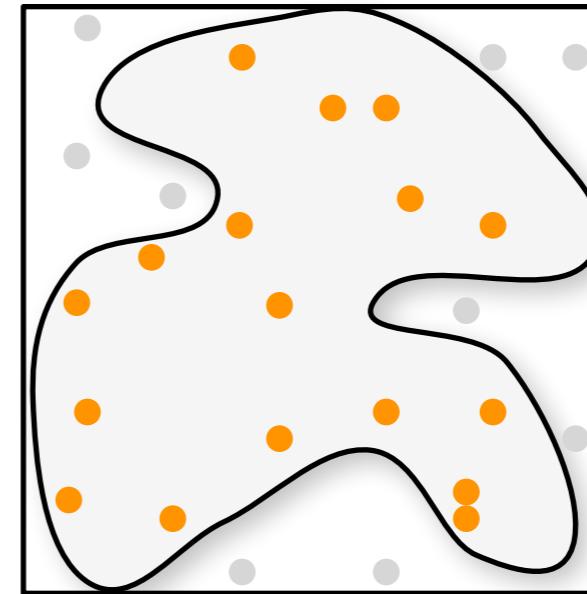
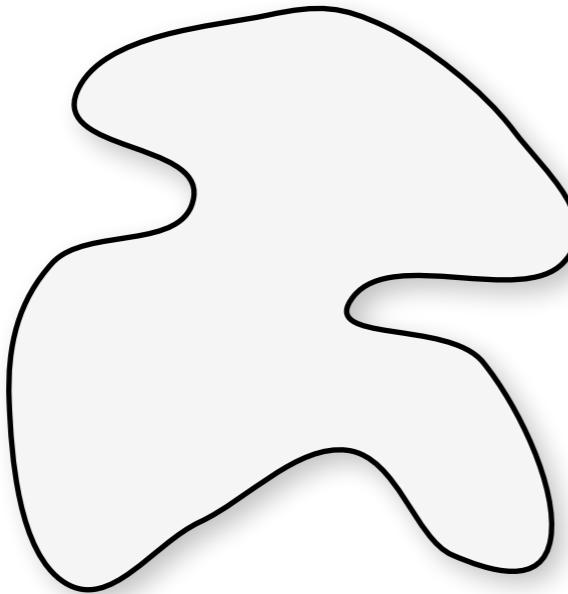


```
Vector3D v;  
do  
{  
    v.x = 1-2*randf();  
    v.y = 1-2*randf();  
    v.z = 1-2*randf();  
} while(v.length2() > 1)  
// flip to proper hemisphere  
if (dot(v,n) < 0)  
    v = -v;  
// Project onto hemisphere  
v /= v.length();
```

- Or, generate in canonical orientation, and then rotate

Rejection Sampling

- More complex shapes



- Pros:
 - + Flexible
- Cons:
 - Inefficient
 - Difficult to combine with stratification or quasi-Monte Carlo

Sampling Arbitrary Distributions

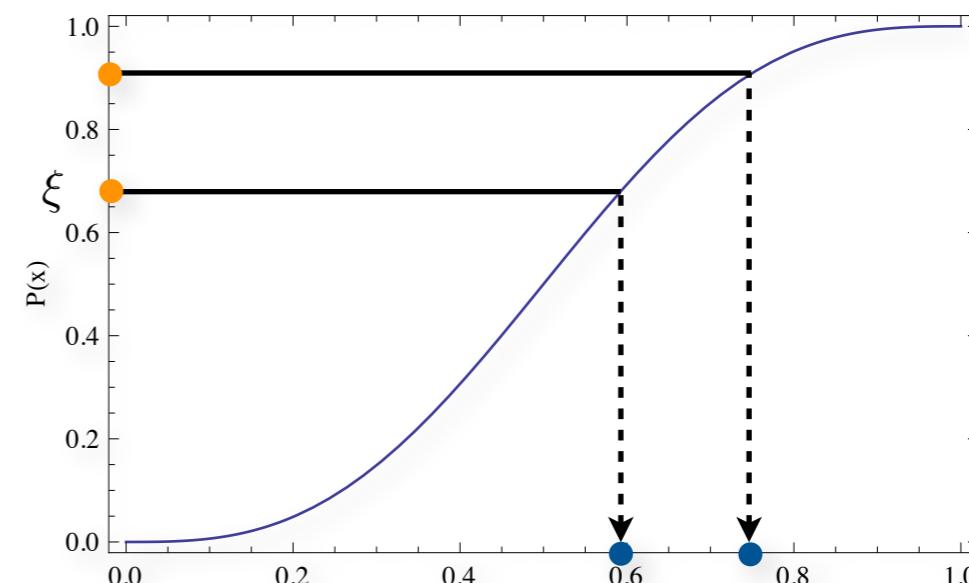
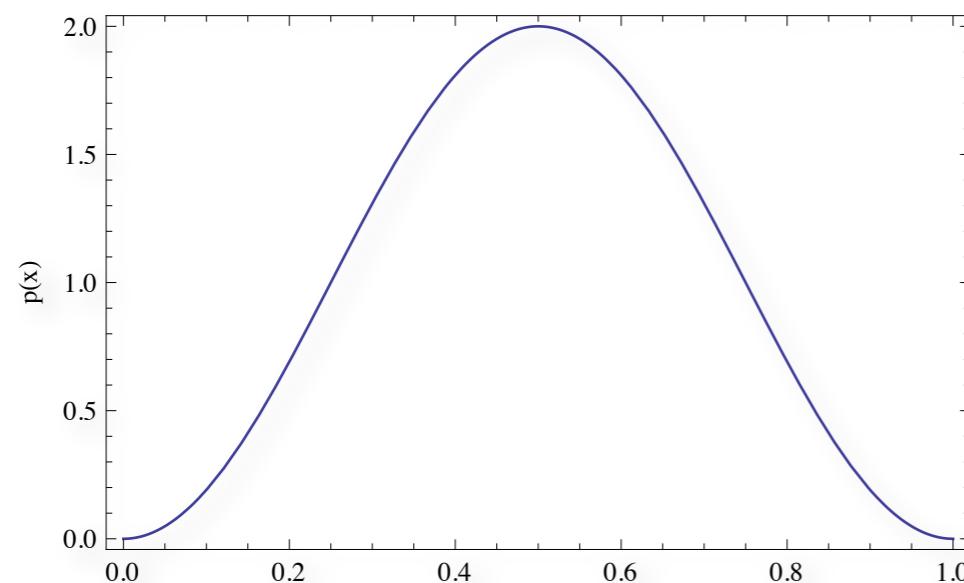
- The inversion method:

1. Compute the CDF $P(x) = \int_0^x p(x') dx'$

2. Compute its inverse $P^{-1}(x)$

3. Obtain a uniformly distributed random number ξ

4. Compute $X_i = P^{-1}(\xi)$

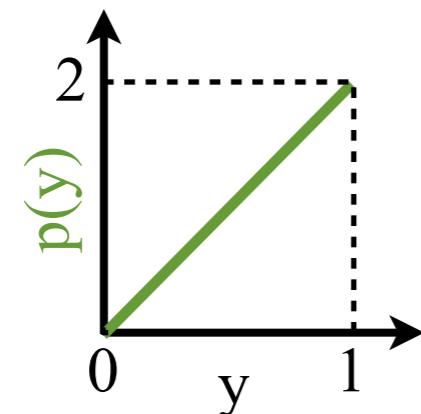


Example – Sampling a Linear Ramp

- Goal: sample with PDF

$$p(y) = 2y$$

$$P(y) = y^2$$



$$P^{-1}(y) = \sqrt{y}$$

$$Y_i = \sqrt{\xi}$$

Sampling 2D Distributions

- Draw samples (X, Y) from a 2D distribution $p(x, y)$
- If $p(x, y)$ is separable, i.e., $p(x, y) = p_x(x) p_y(y)$, we can independently sample $p_x(x)$, and $p_y(y)$
- Otherwise, compute the marginal density function:
$$p(x) = \int p(x, y) dy$$
- and, the conditional density:
$$p(y|x) = \frac{p(x, y)}{p(x)}$$
- Procedure: first sample $X_i \sim p(x)$, then $Y_i \sim p(y|x)$

Visual Break



Andreas Byström